



人工智能技术及应用


Artificial Intelligence and Application

Classification




Classification




- Credit Scoring
 - Input: income, savings, profession, age, past financial history
 - Output: accept or refuse
- Medical Diagnosis
 - Input: current symptoms, age, gender, past medical history
 - Output: which kind of diseases
- Handwritten character recognition
 - Input:  output: 金
- Face recognition
 - Input: image of a face, output: person

Classification as Regression?



- Regression



- Classification as regression?



different? $\left[\begin{array}{l} \rightarrow 1 = \text{class 1} \\ \rightarrow 2 = \text{class 2} \\ \rightarrow 3 = \text{class 3} \end{array} \right. \left. \begin{array}{l} \leftarrow \\ \leftarrow \\ \leftarrow \end{array} \right] \text{ similar?}$



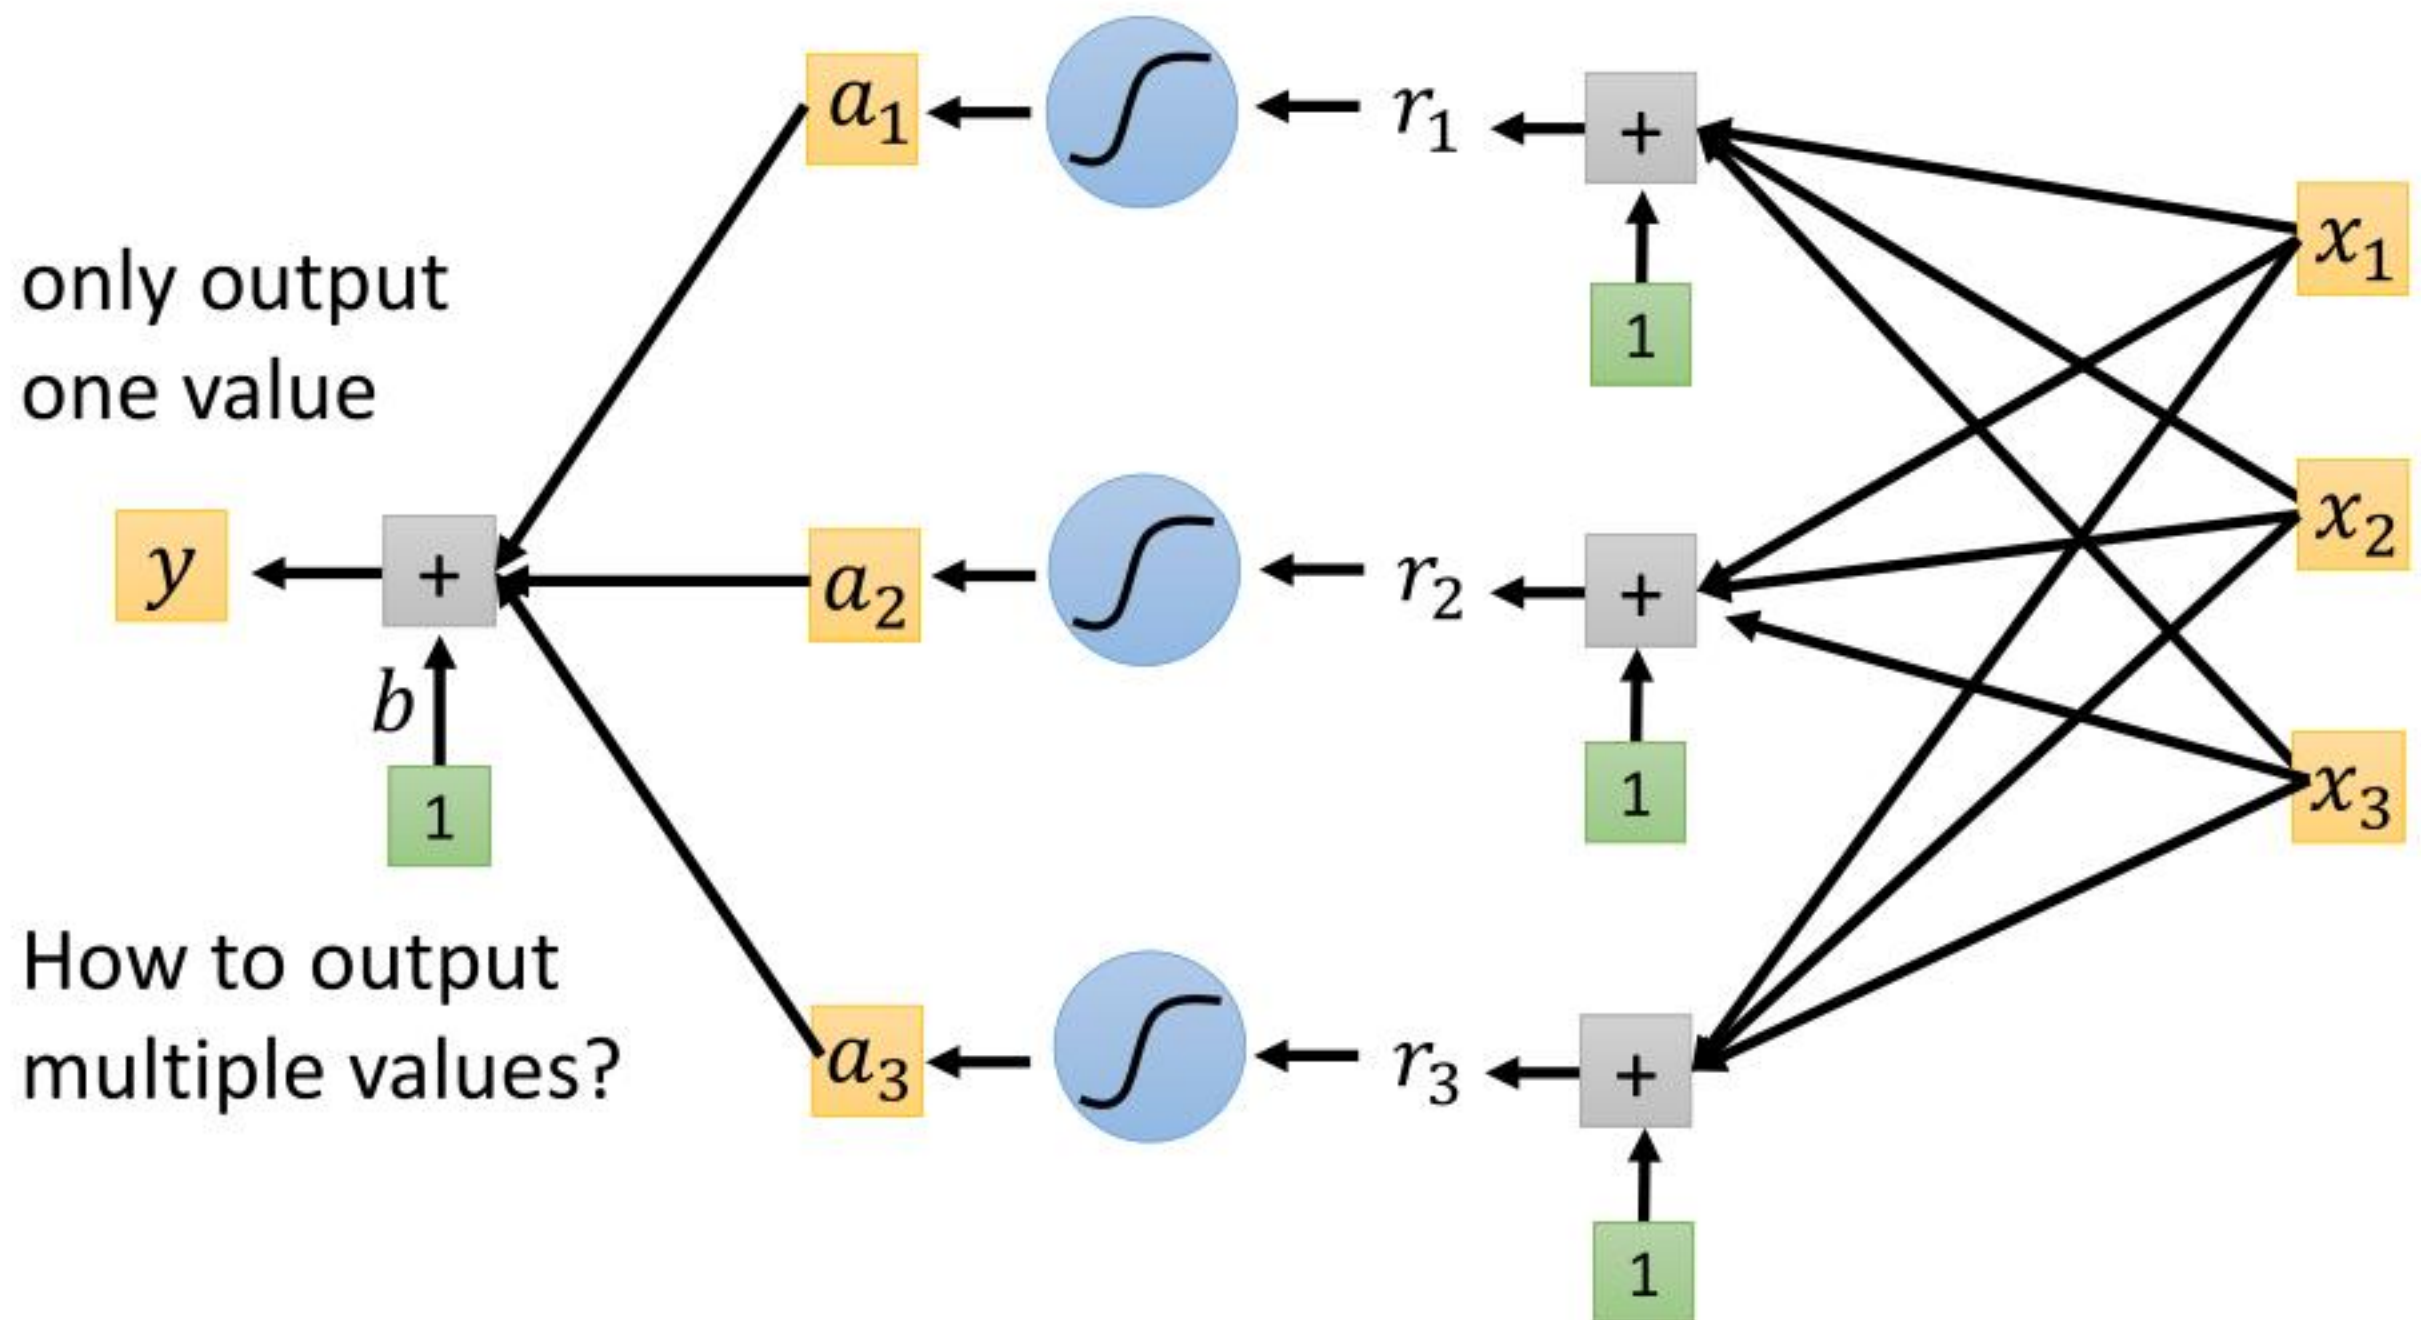
Penalize to the examples that are “too correct” ... (Bishop, P186)

- Multiple class: Class 1 means the target is 1; Class 2 means the target is 2; Class 3 means the target is 3 problematic

Class as one-hot vector

Class 1 Class 2 Class 3


$$\hat{\mathbf{y}} = \begin{bmatrix} 1 \\ 0 \\ 0 \end{bmatrix} \quad \text{or} \quad \begin{bmatrix} 0 \\ 1 \\ 0 \end{bmatrix} \quad \text{or} \quad \begin{bmatrix} 0 \\ 0 \\ 1 \end{bmatrix}$$



Class as one-hot vector

Class 1 Class 2 Class 3

$$\hat{\mathbf{y}} = \begin{bmatrix} 1 \\ 0 \\ 0 \end{bmatrix} \quad \text{or} \quad \begin{bmatrix} 0 \\ 1 \\ 0 \end{bmatrix} \quad \text{or} \quad \begin{bmatrix} 0 \\ 0 \\ 1 \end{bmatrix}$$

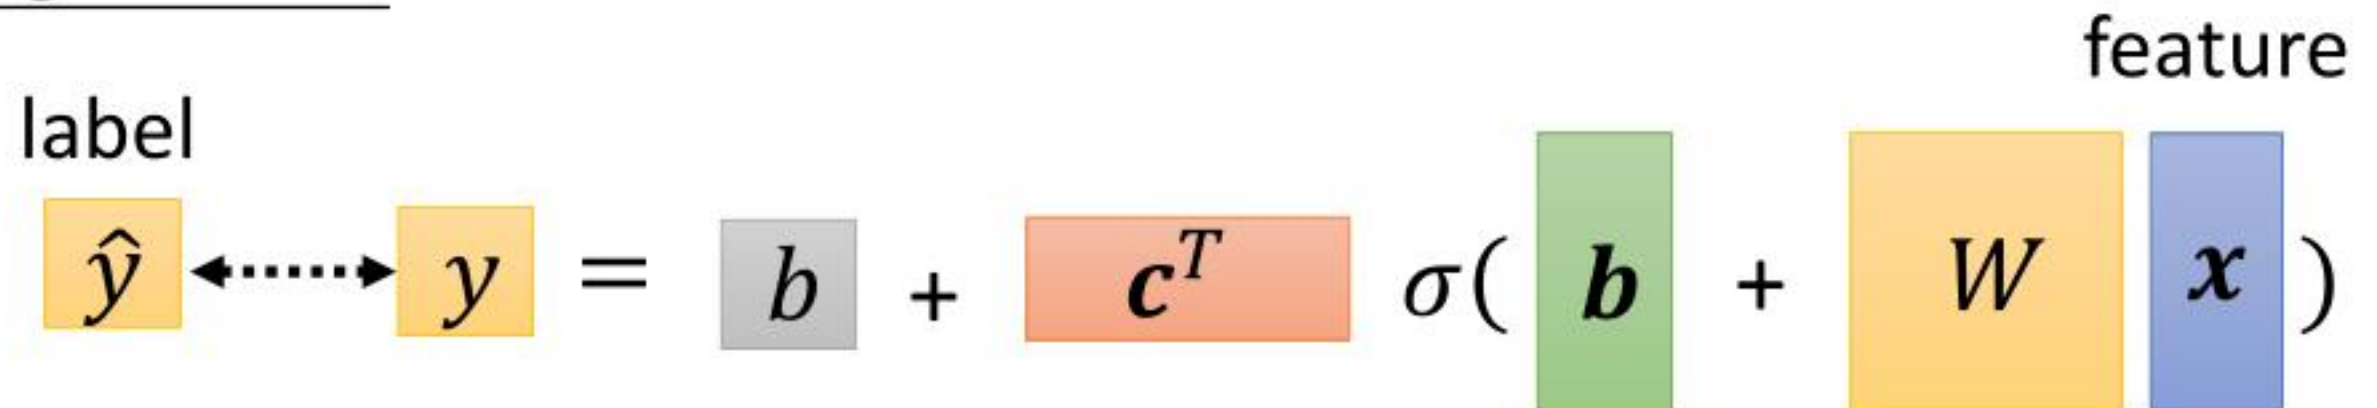


Regression

label

$$\hat{y} \longleftrightarrow y = b + c^T \sigma(b + W x)$$

feature



Classification

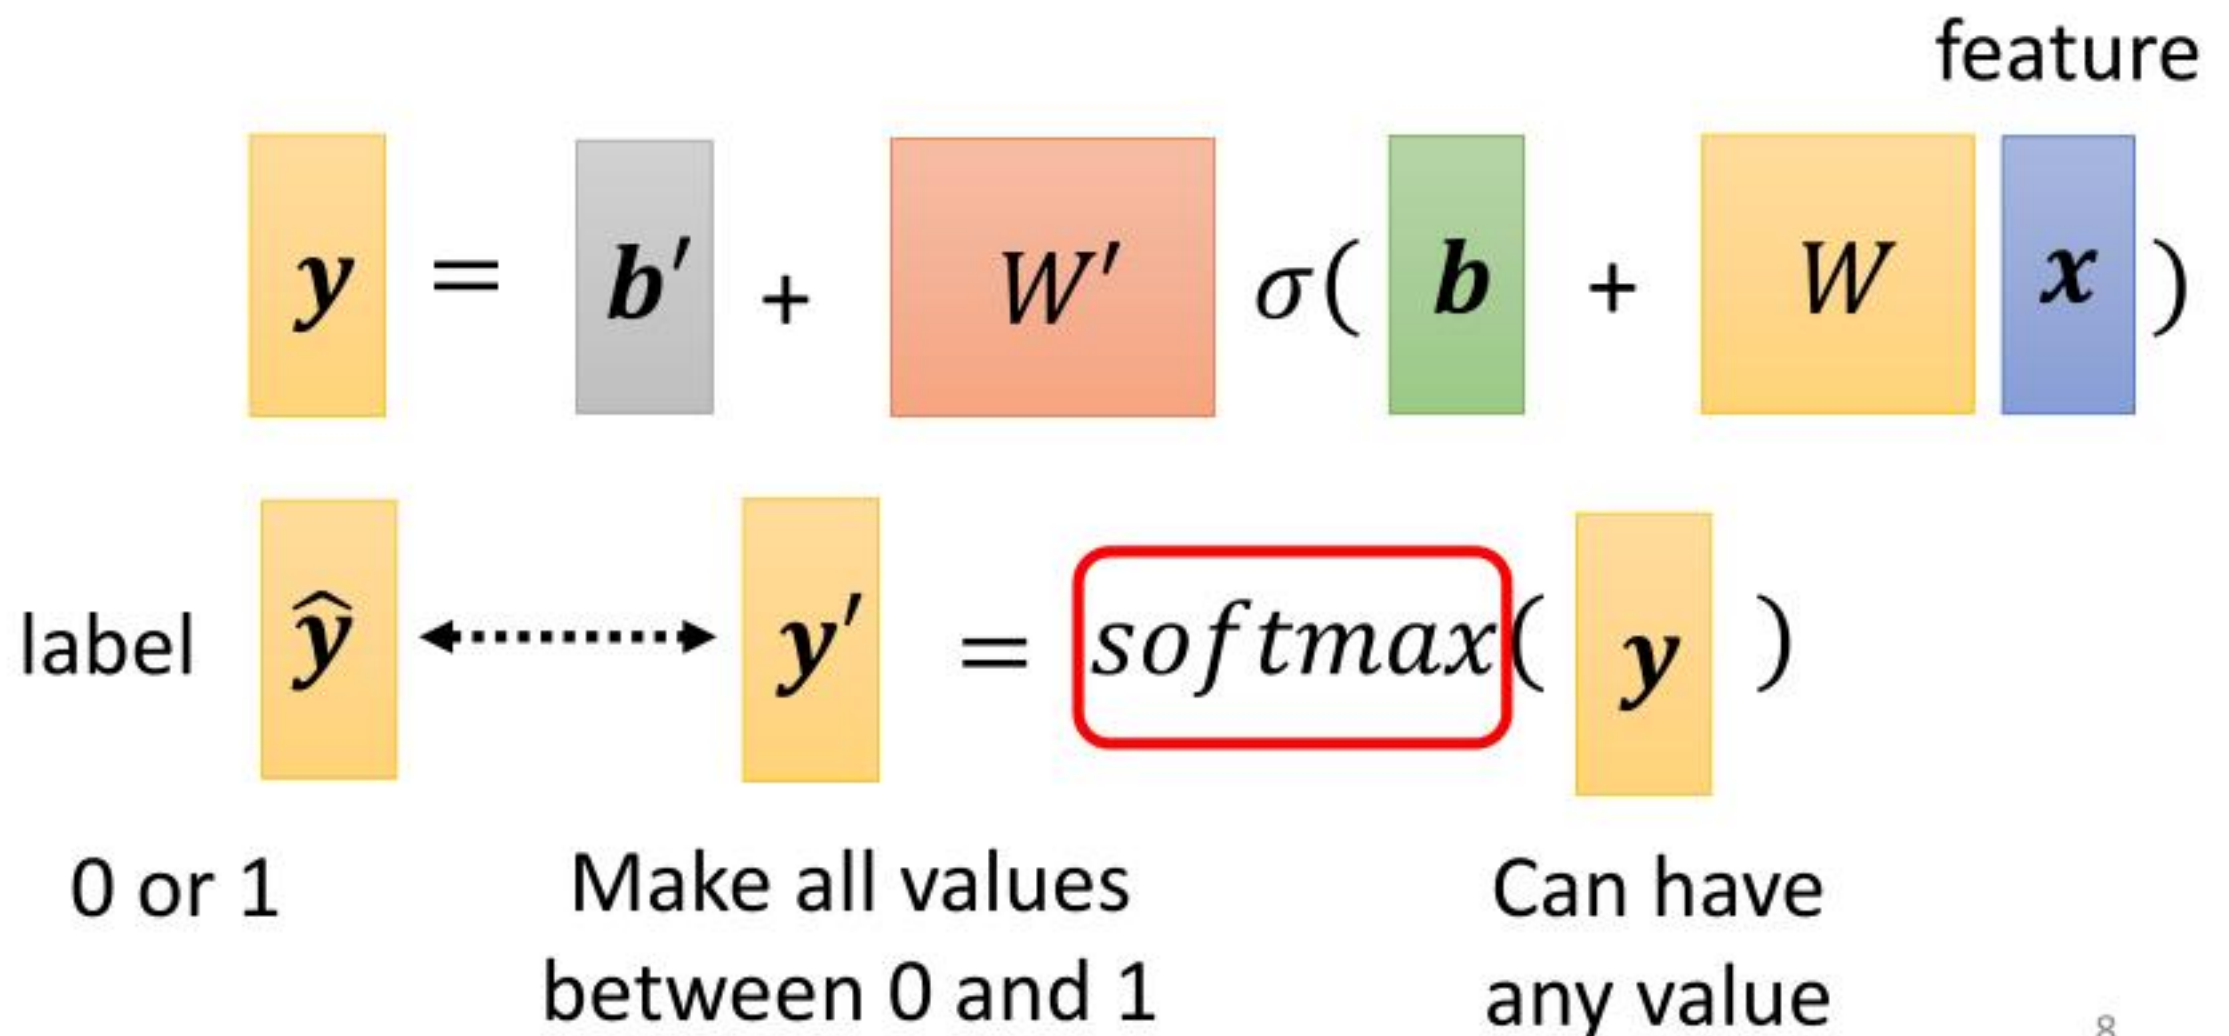
feature

$$y = b' + W' \sigma(b + W x)$$

label

$$\hat{y} \longleftrightarrow y' = \text{softmax}(y)$$

0 or 1 Make all values between 0 and 1 Can have any value




Soft-max


$$y'_i = \frac{\exp(y_i)}{\sum_j \exp(y_i)} \quad \begin{array}{l} \blacksquare 1 > y'_i > 0 \\ \blacksquare \sum_i y'_i = 1 \end{array}$$

Softmax

How about **binary classification**? 😊



Multi-class Classification (3 classes as example)



If $x \in \text{class 1}$

$$\hat{y} = \begin{bmatrix} 1 \\ 0 \\ 0 \end{bmatrix}$$

$$-\ln y_1$$

If $x \in \text{class 2}$

$$\hat{y} = \begin{bmatrix} 0 \\ 1 \\ 0 \end{bmatrix}$$

$$-\ln y_2$$


If $x \in \text{class 3}$

$$\hat{y} = \begin{bmatrix} 0 \\ 0 \\ 1 \end{bmatrix}$$

$$-\ln y_3$$

Loss of Classification

$$L = \frac{1}{N} \sum_n e_n$$



Mean Square Error (MSE)

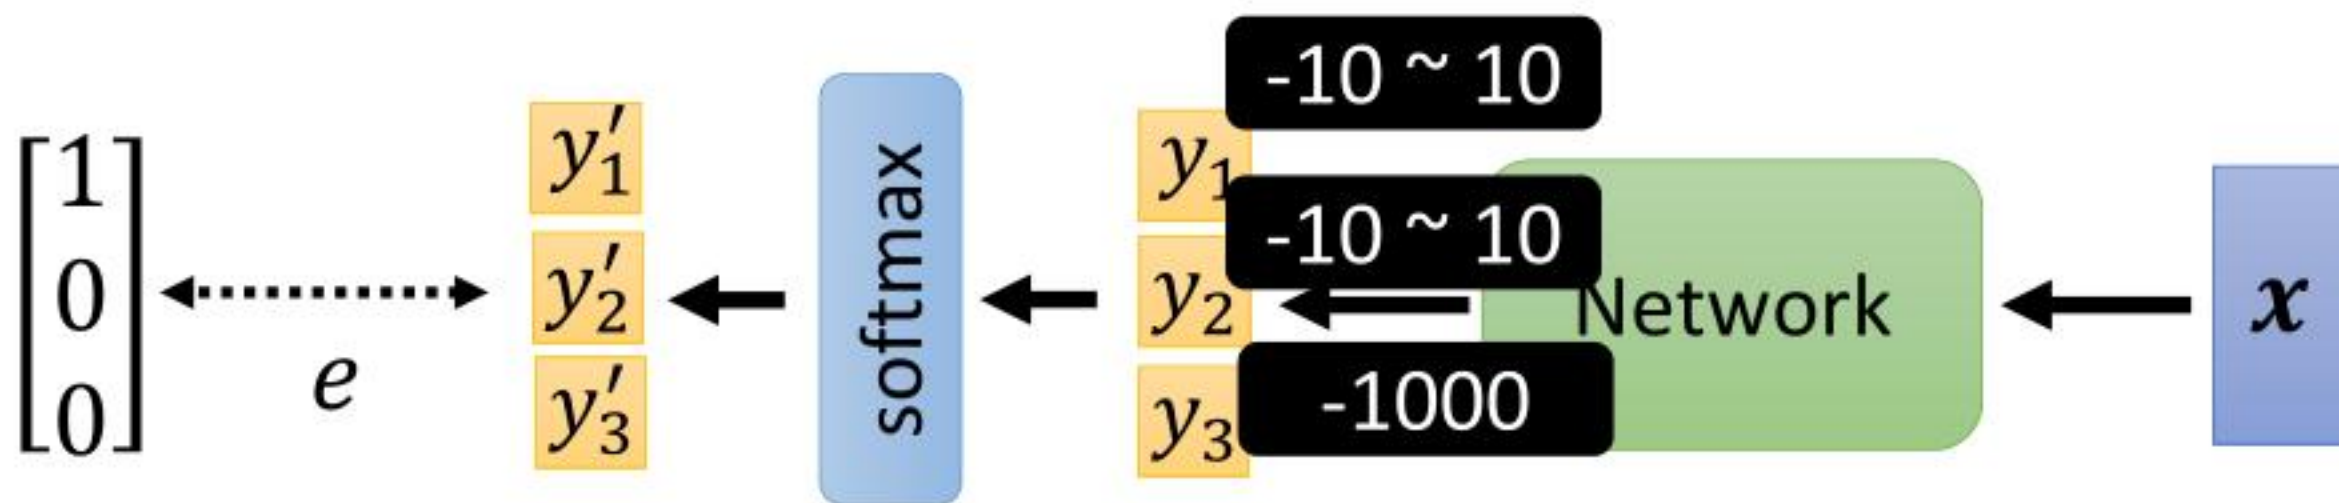
$$e = \sum_i (\hat{y}_i - y'_i)^2$$

Cross-entropy




$$e = - \sum_i \hat{y}_i \ln y'_i$$

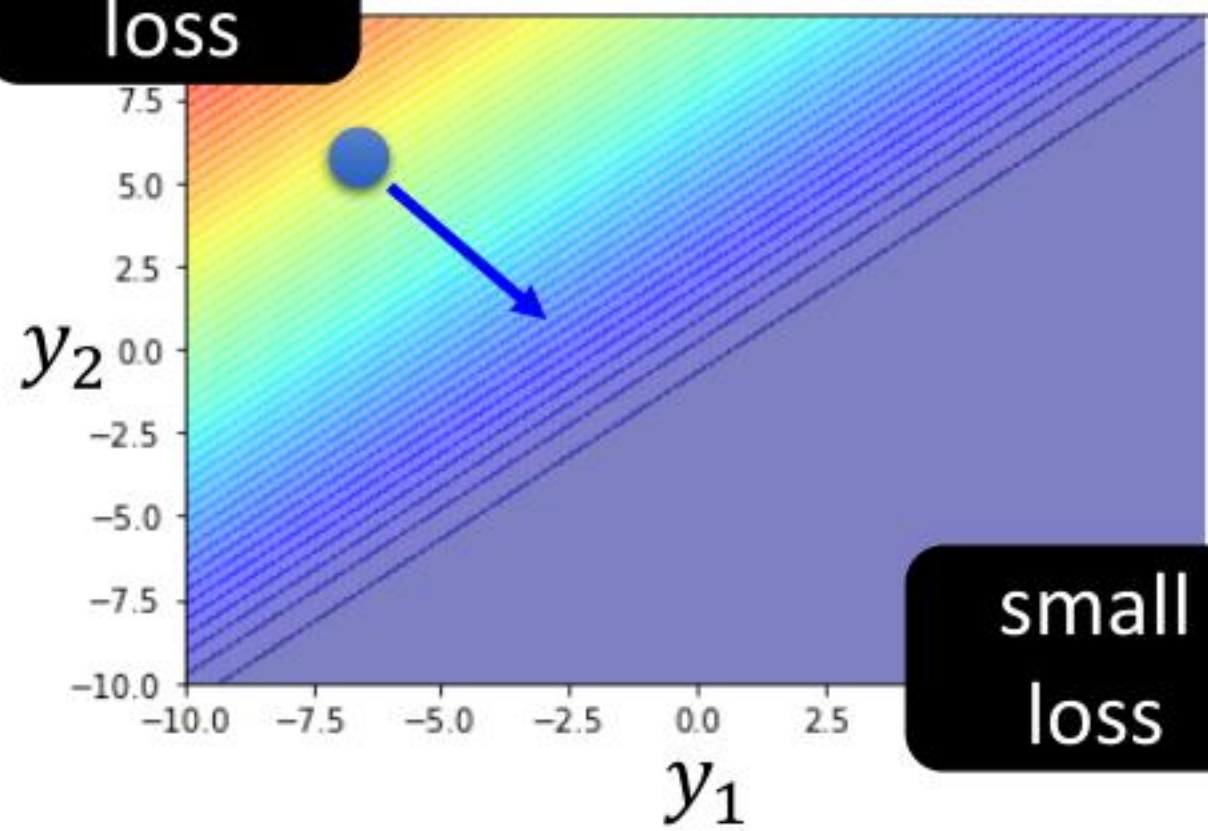
Minimizing cross-entropy is equivalent to **maximizing likelihood**.



large loss Mean Square Error (MSE)



large loss Cross-entropy




Changing the loss function can change the difficulty of optimization.

Convolutional Neural Network (CNN)

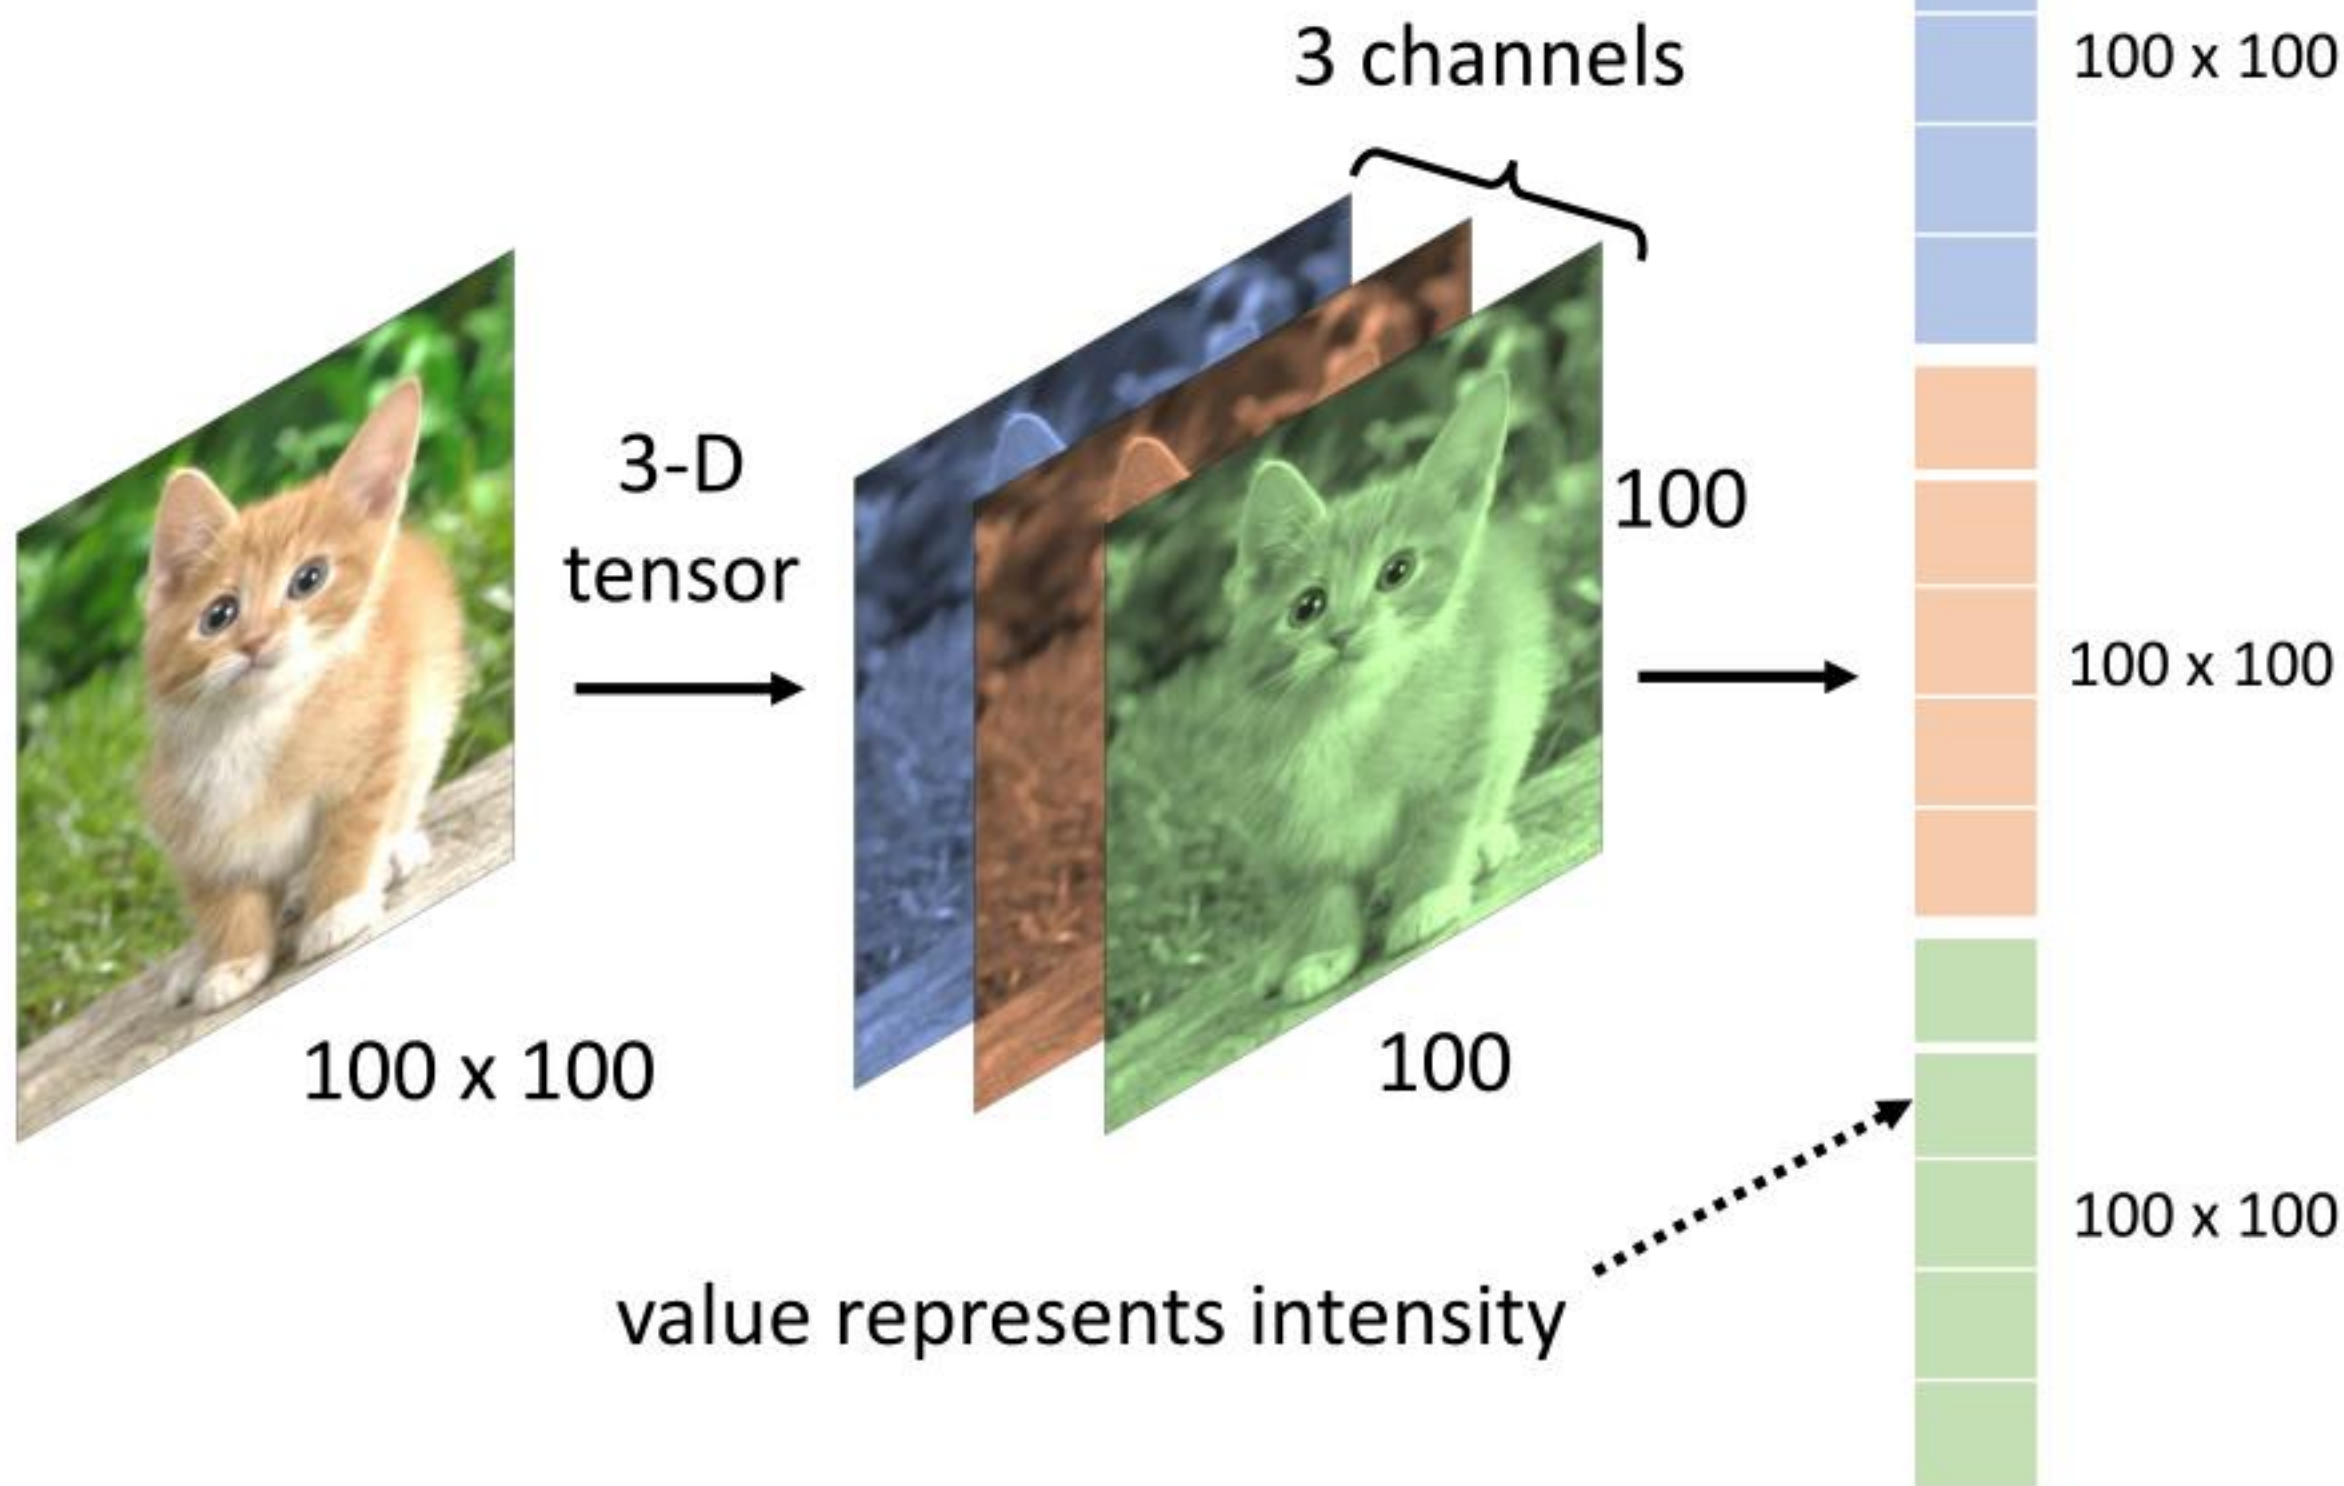


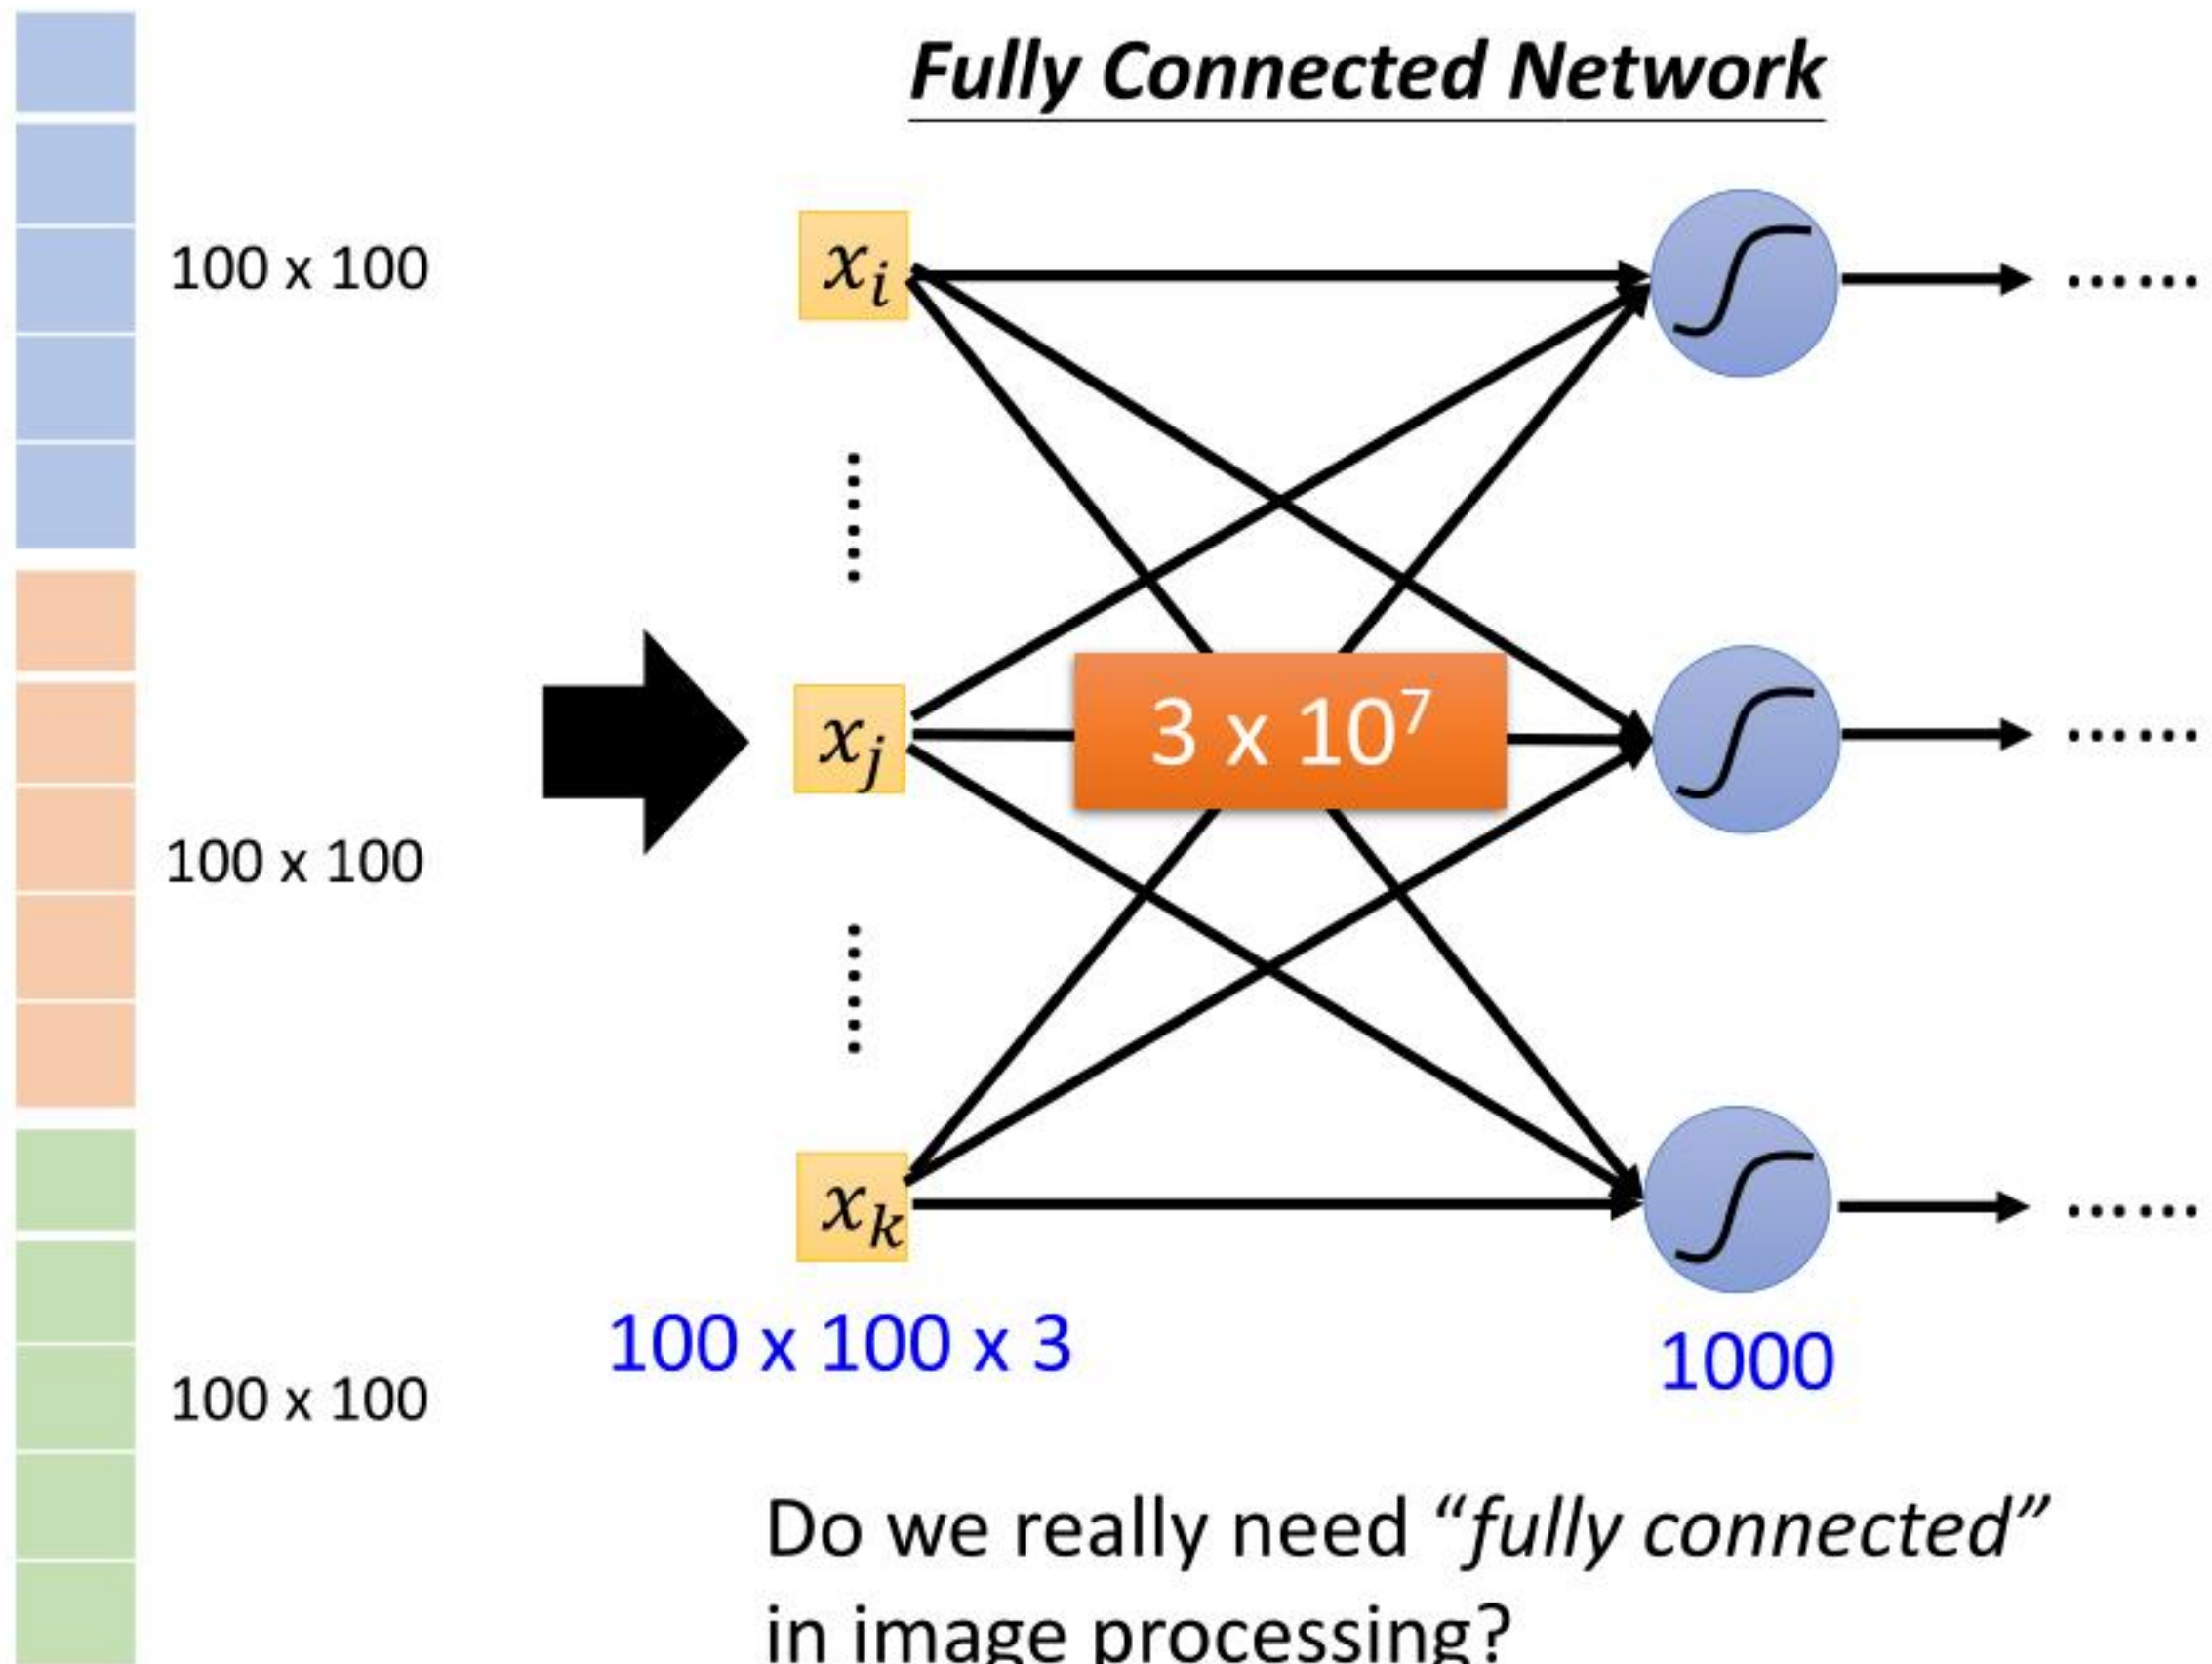
Image Classification



(All the images to be classified have the same size.)

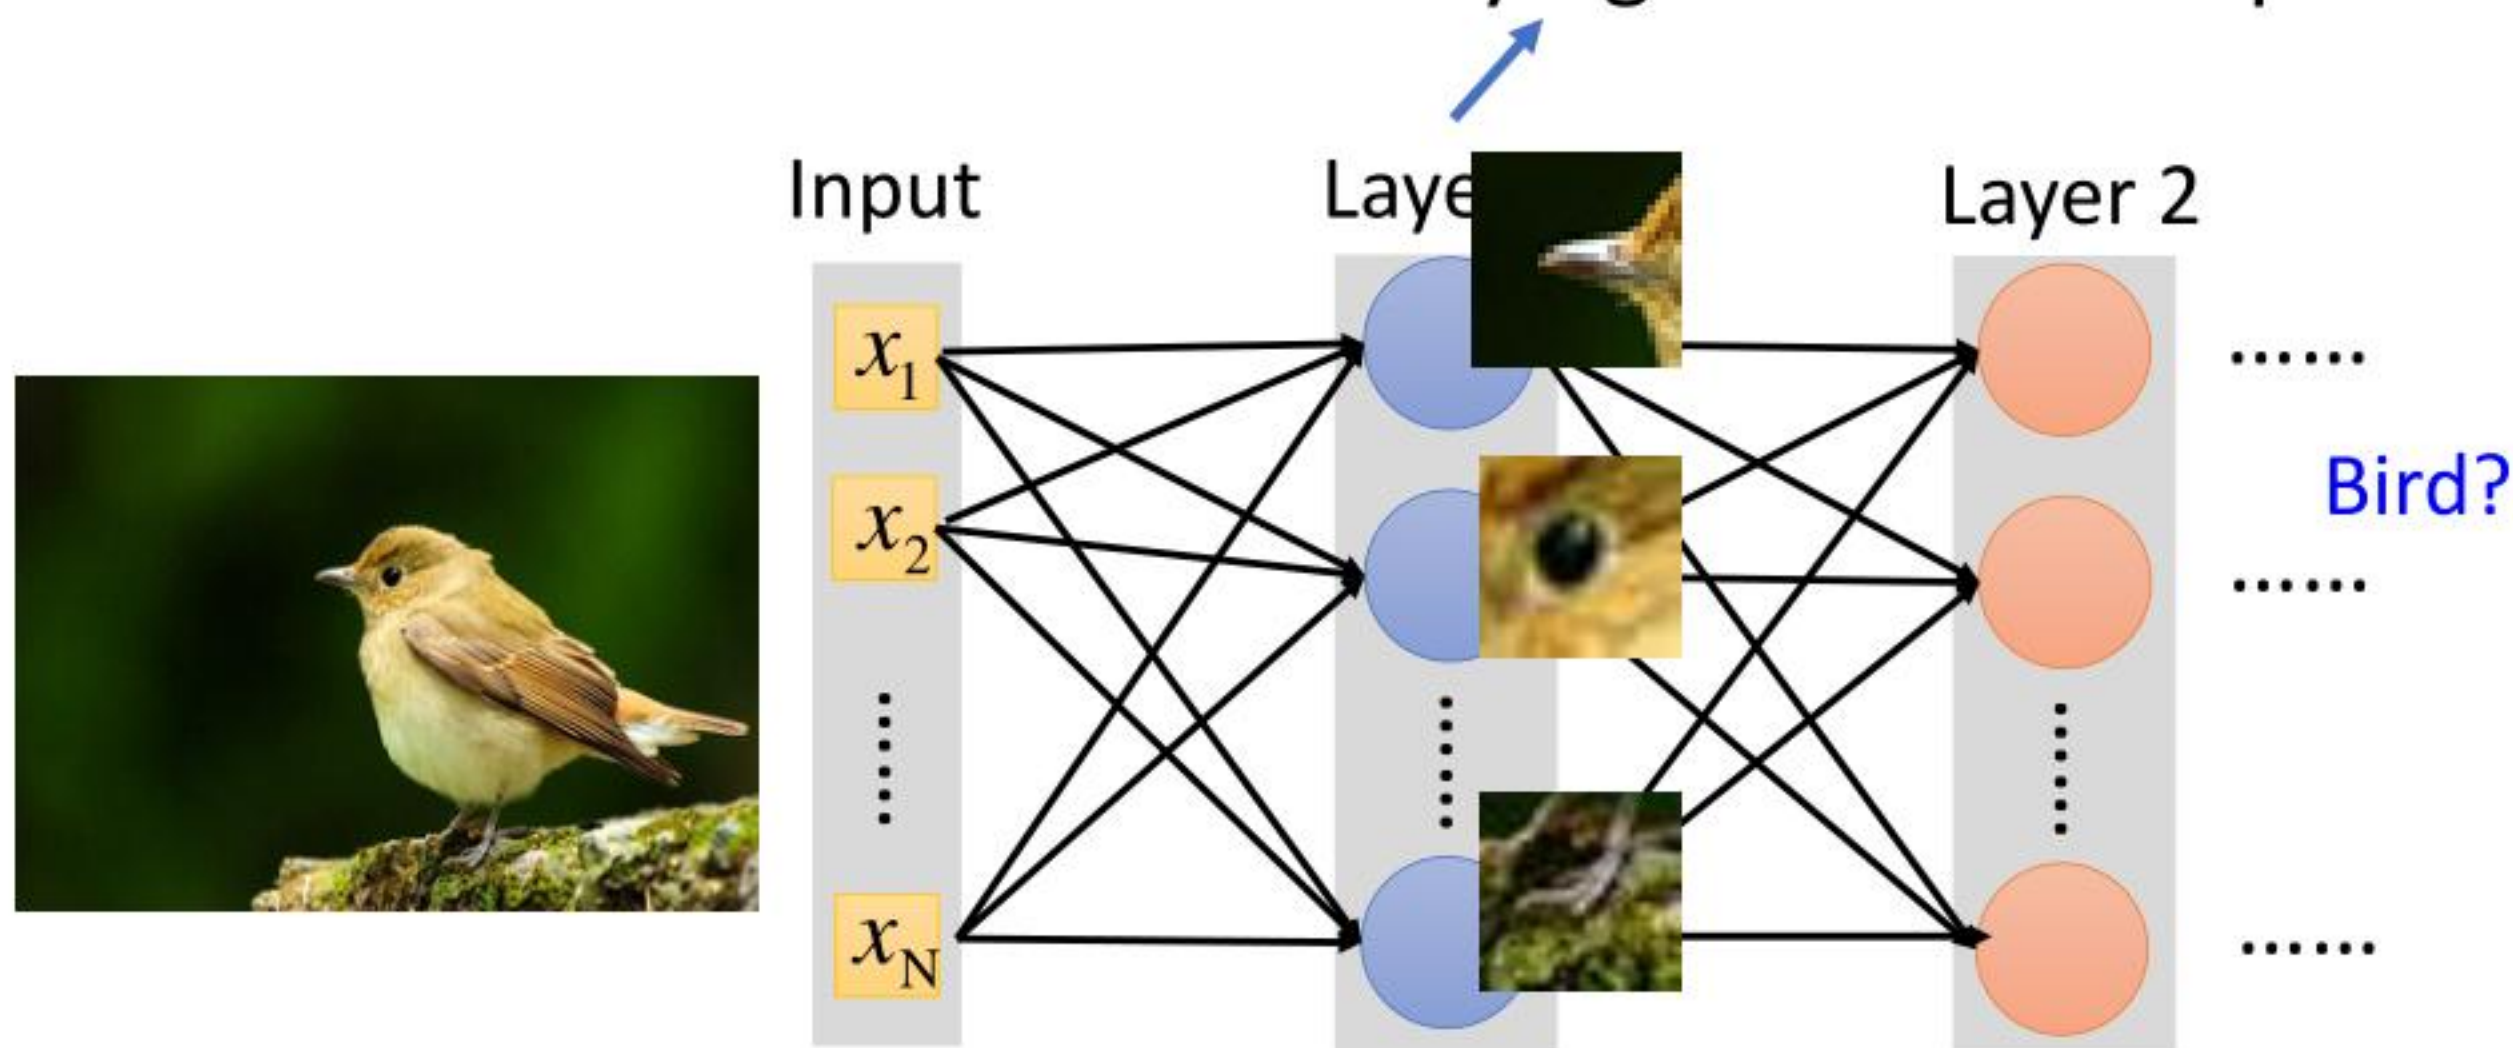
Image Classification





Observation 1

Identifying some critical patterns

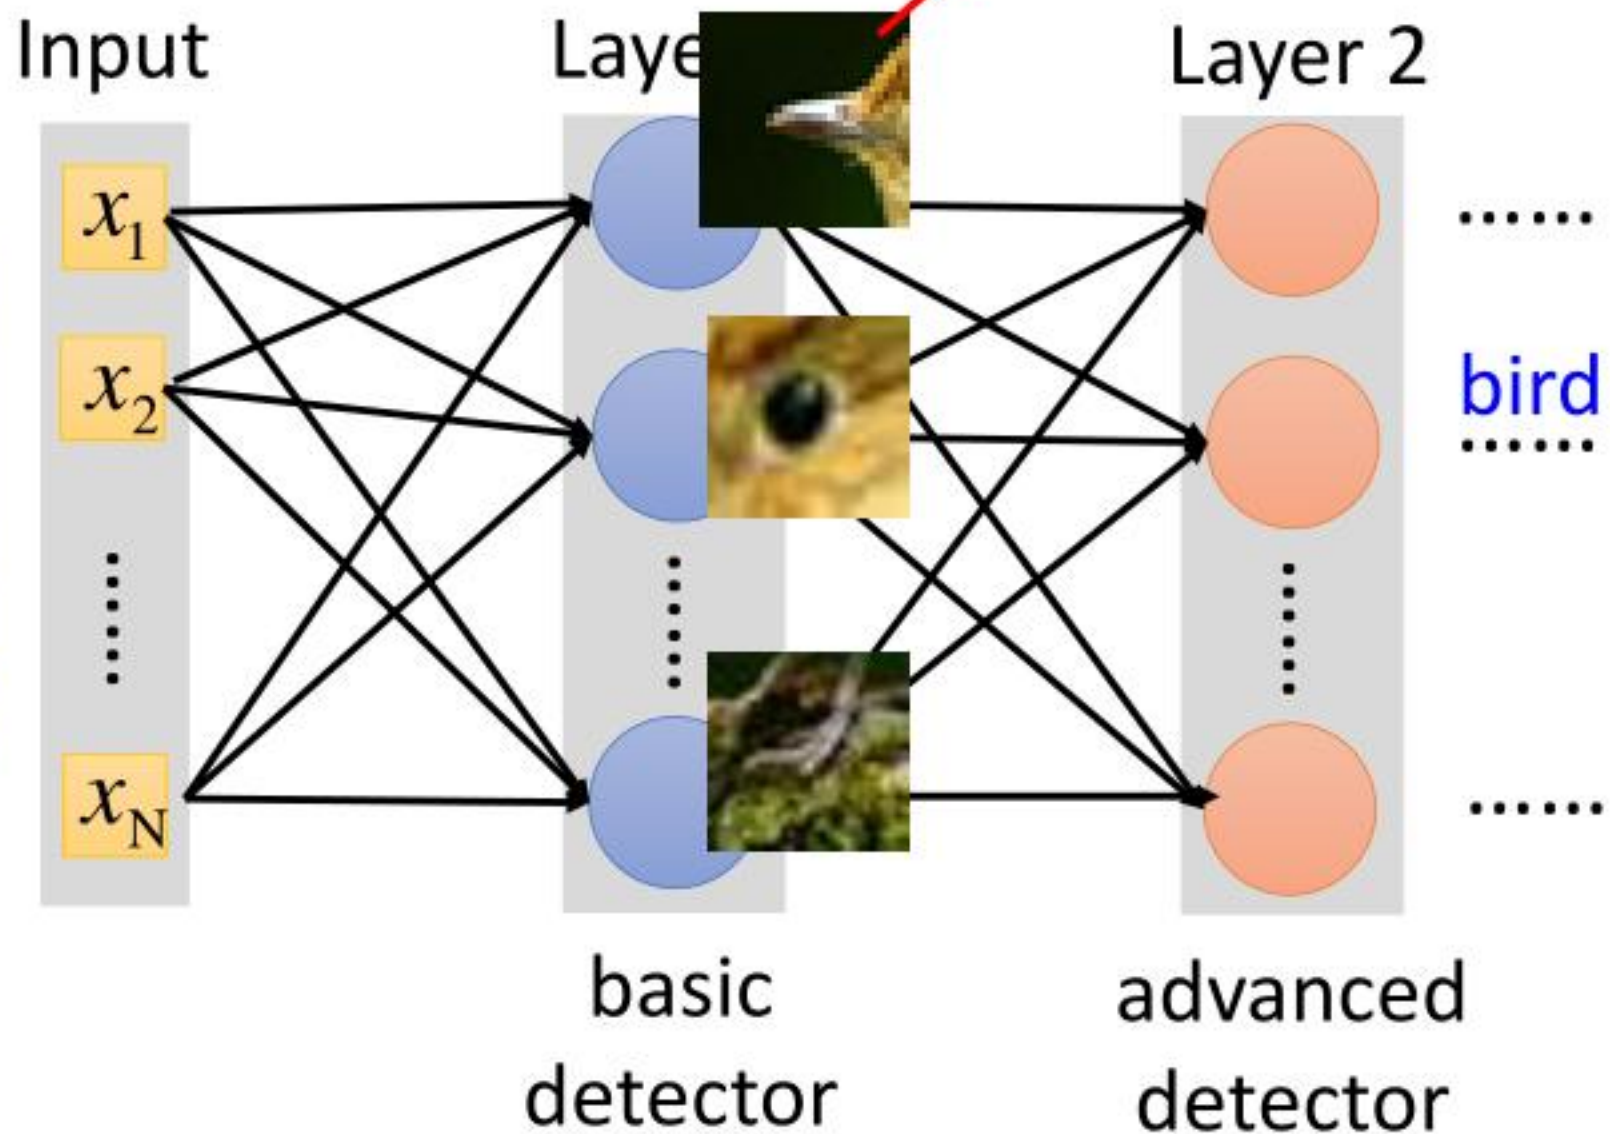



Perhaps human also identify birds in a similar way ... 😊

Observation 1

A neuron does not have to see the whole image.

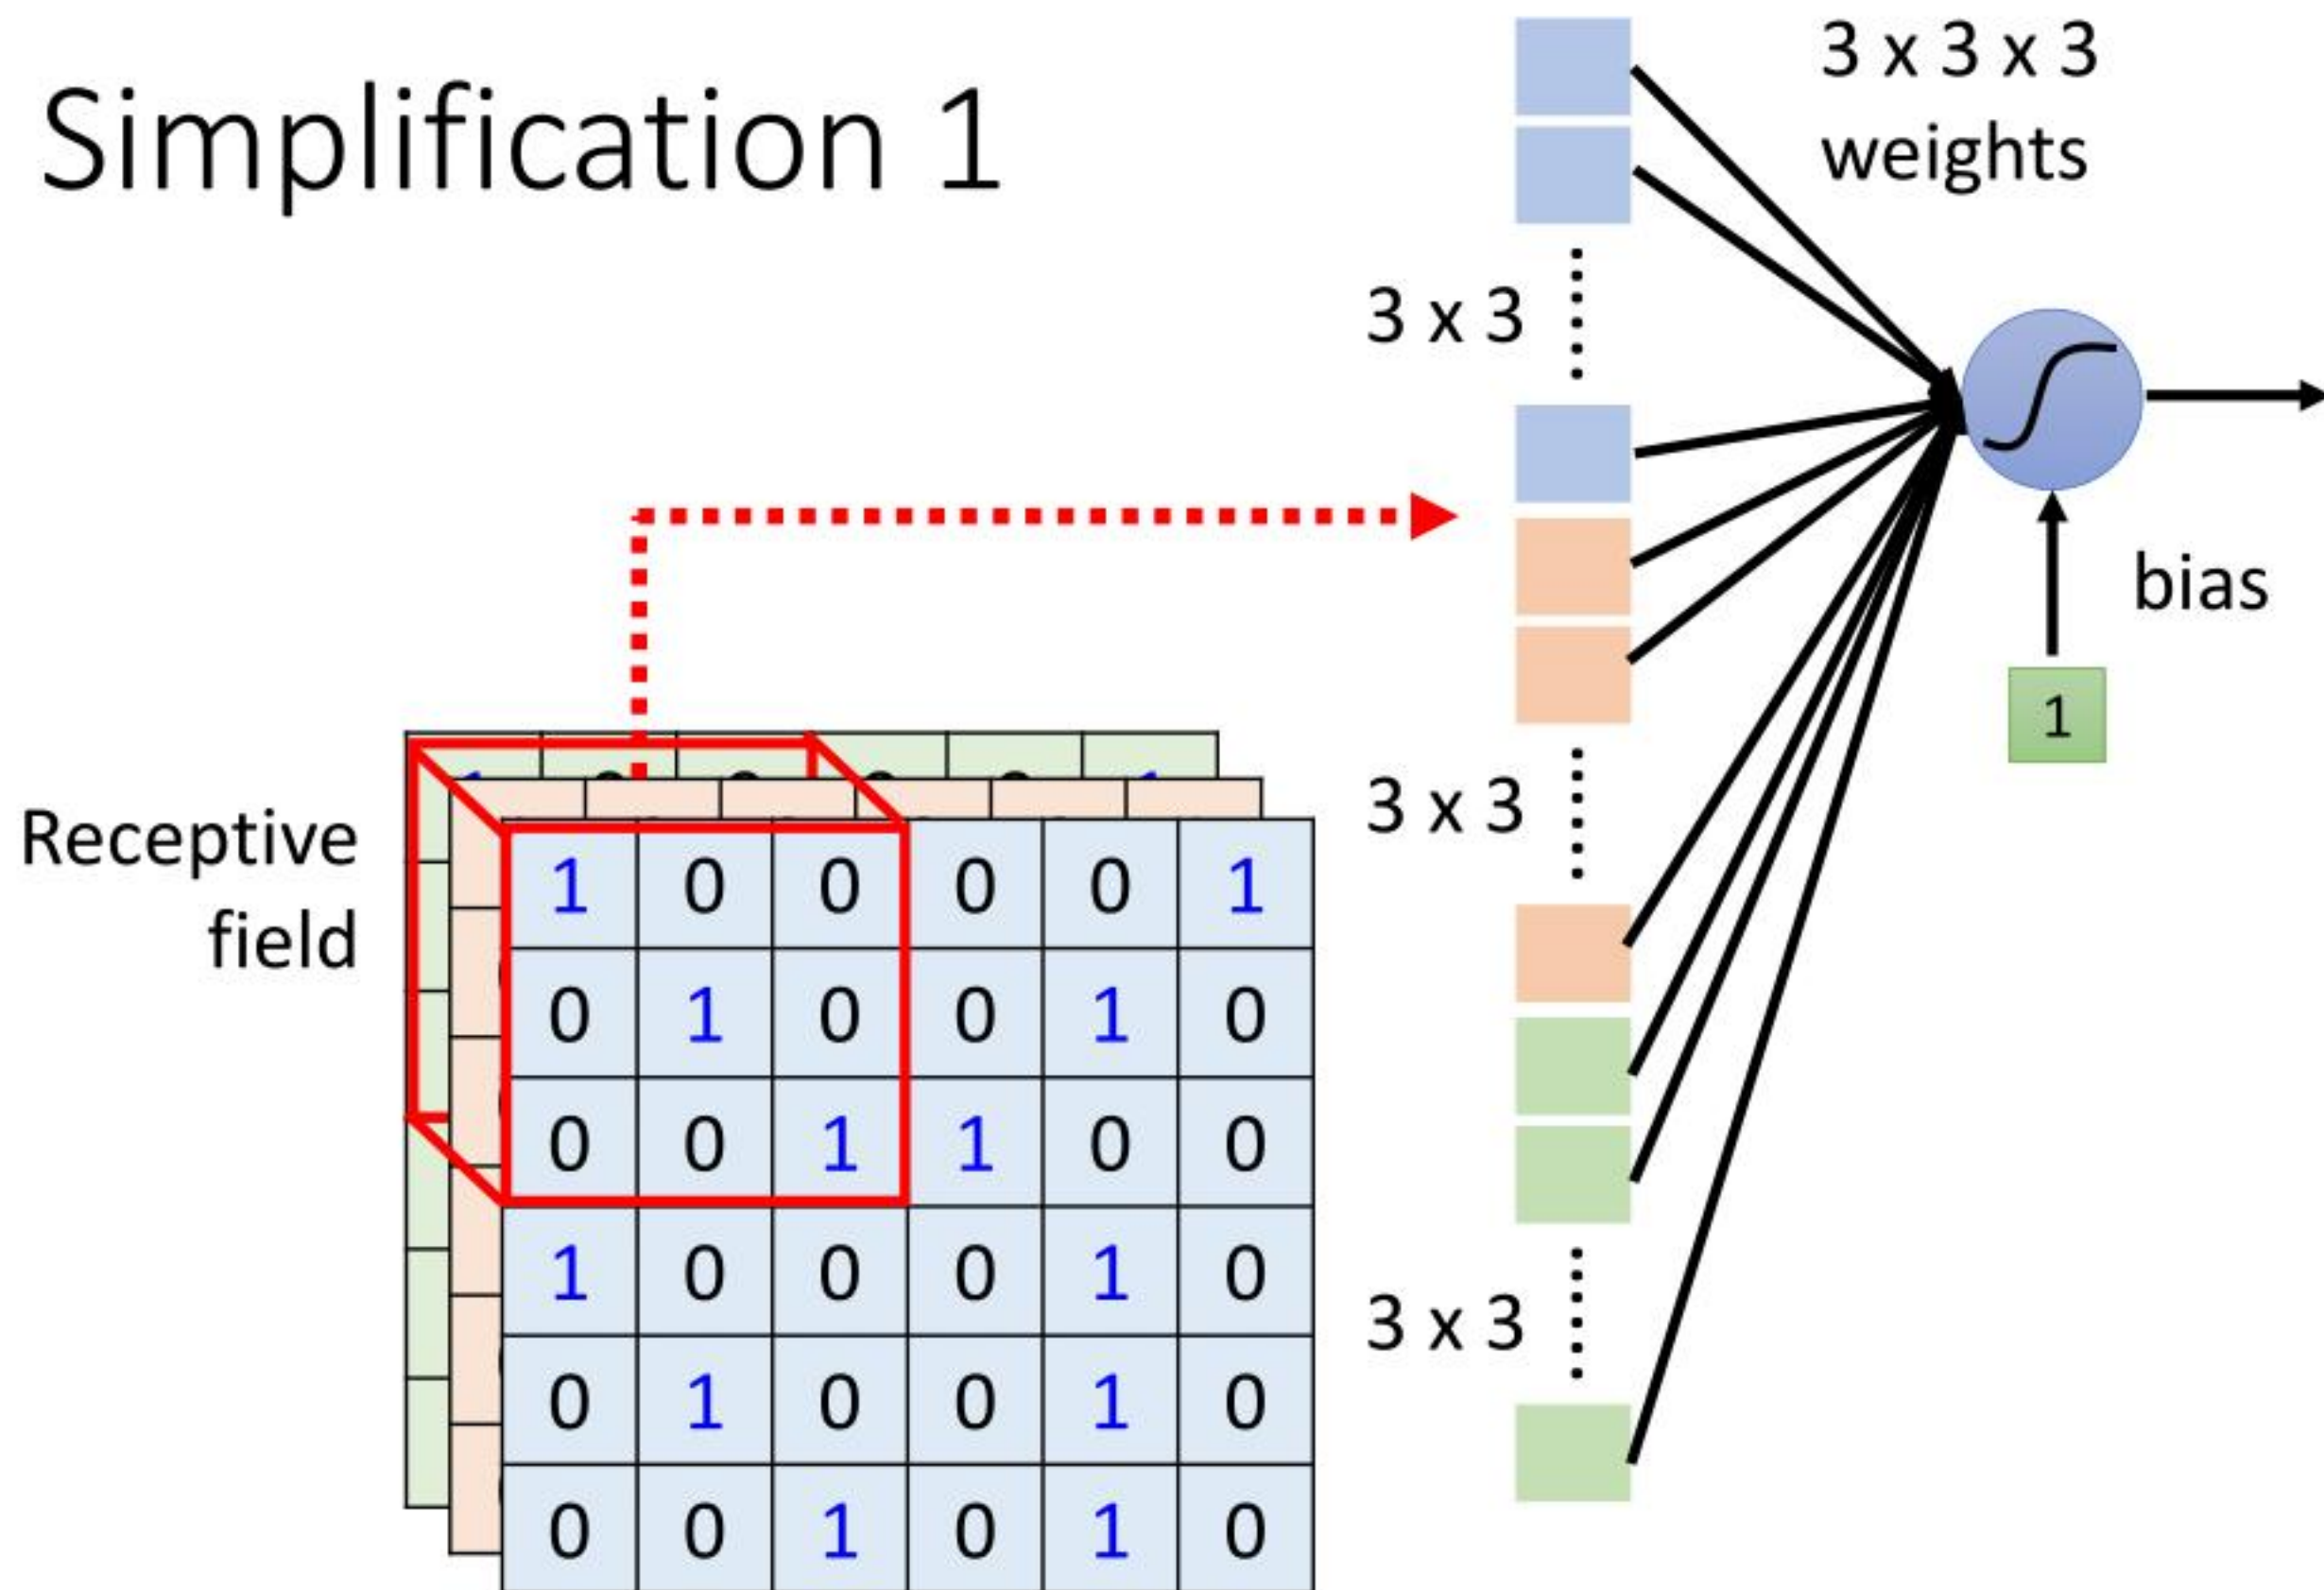
Need to see the whole image?



Some patterns are much smaller than the whole image.


Connecting to small region with less parameters

Simplification 1



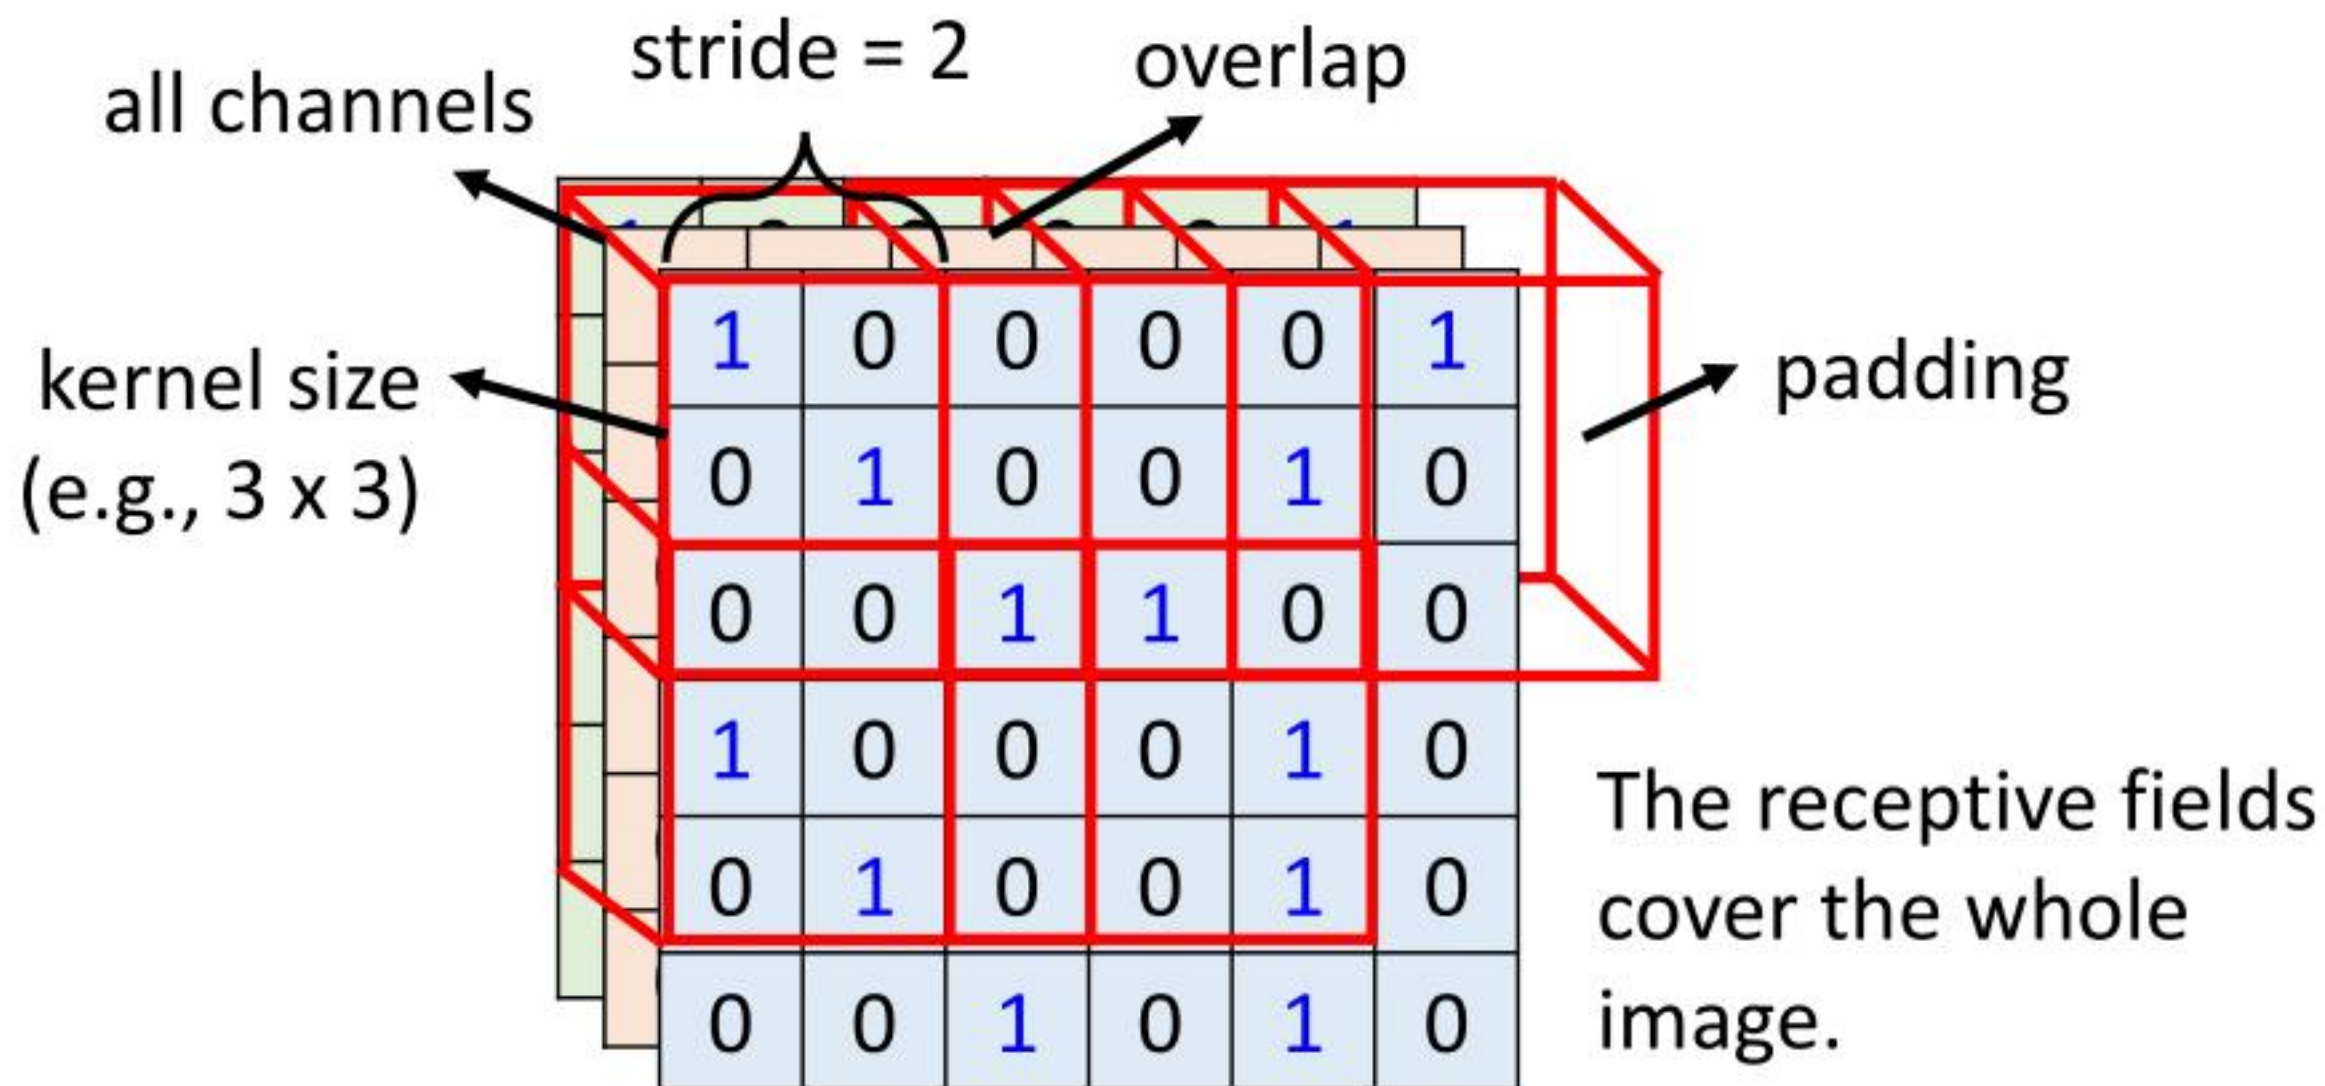
Simplification 1

- Can different neurons have different sizes of receptive field?
- Cover only some channels?
- Not square receptive field?





Simplification 1 – Typical Setting

Each receptive field has a set of neurons (e.g., 64 neurons).

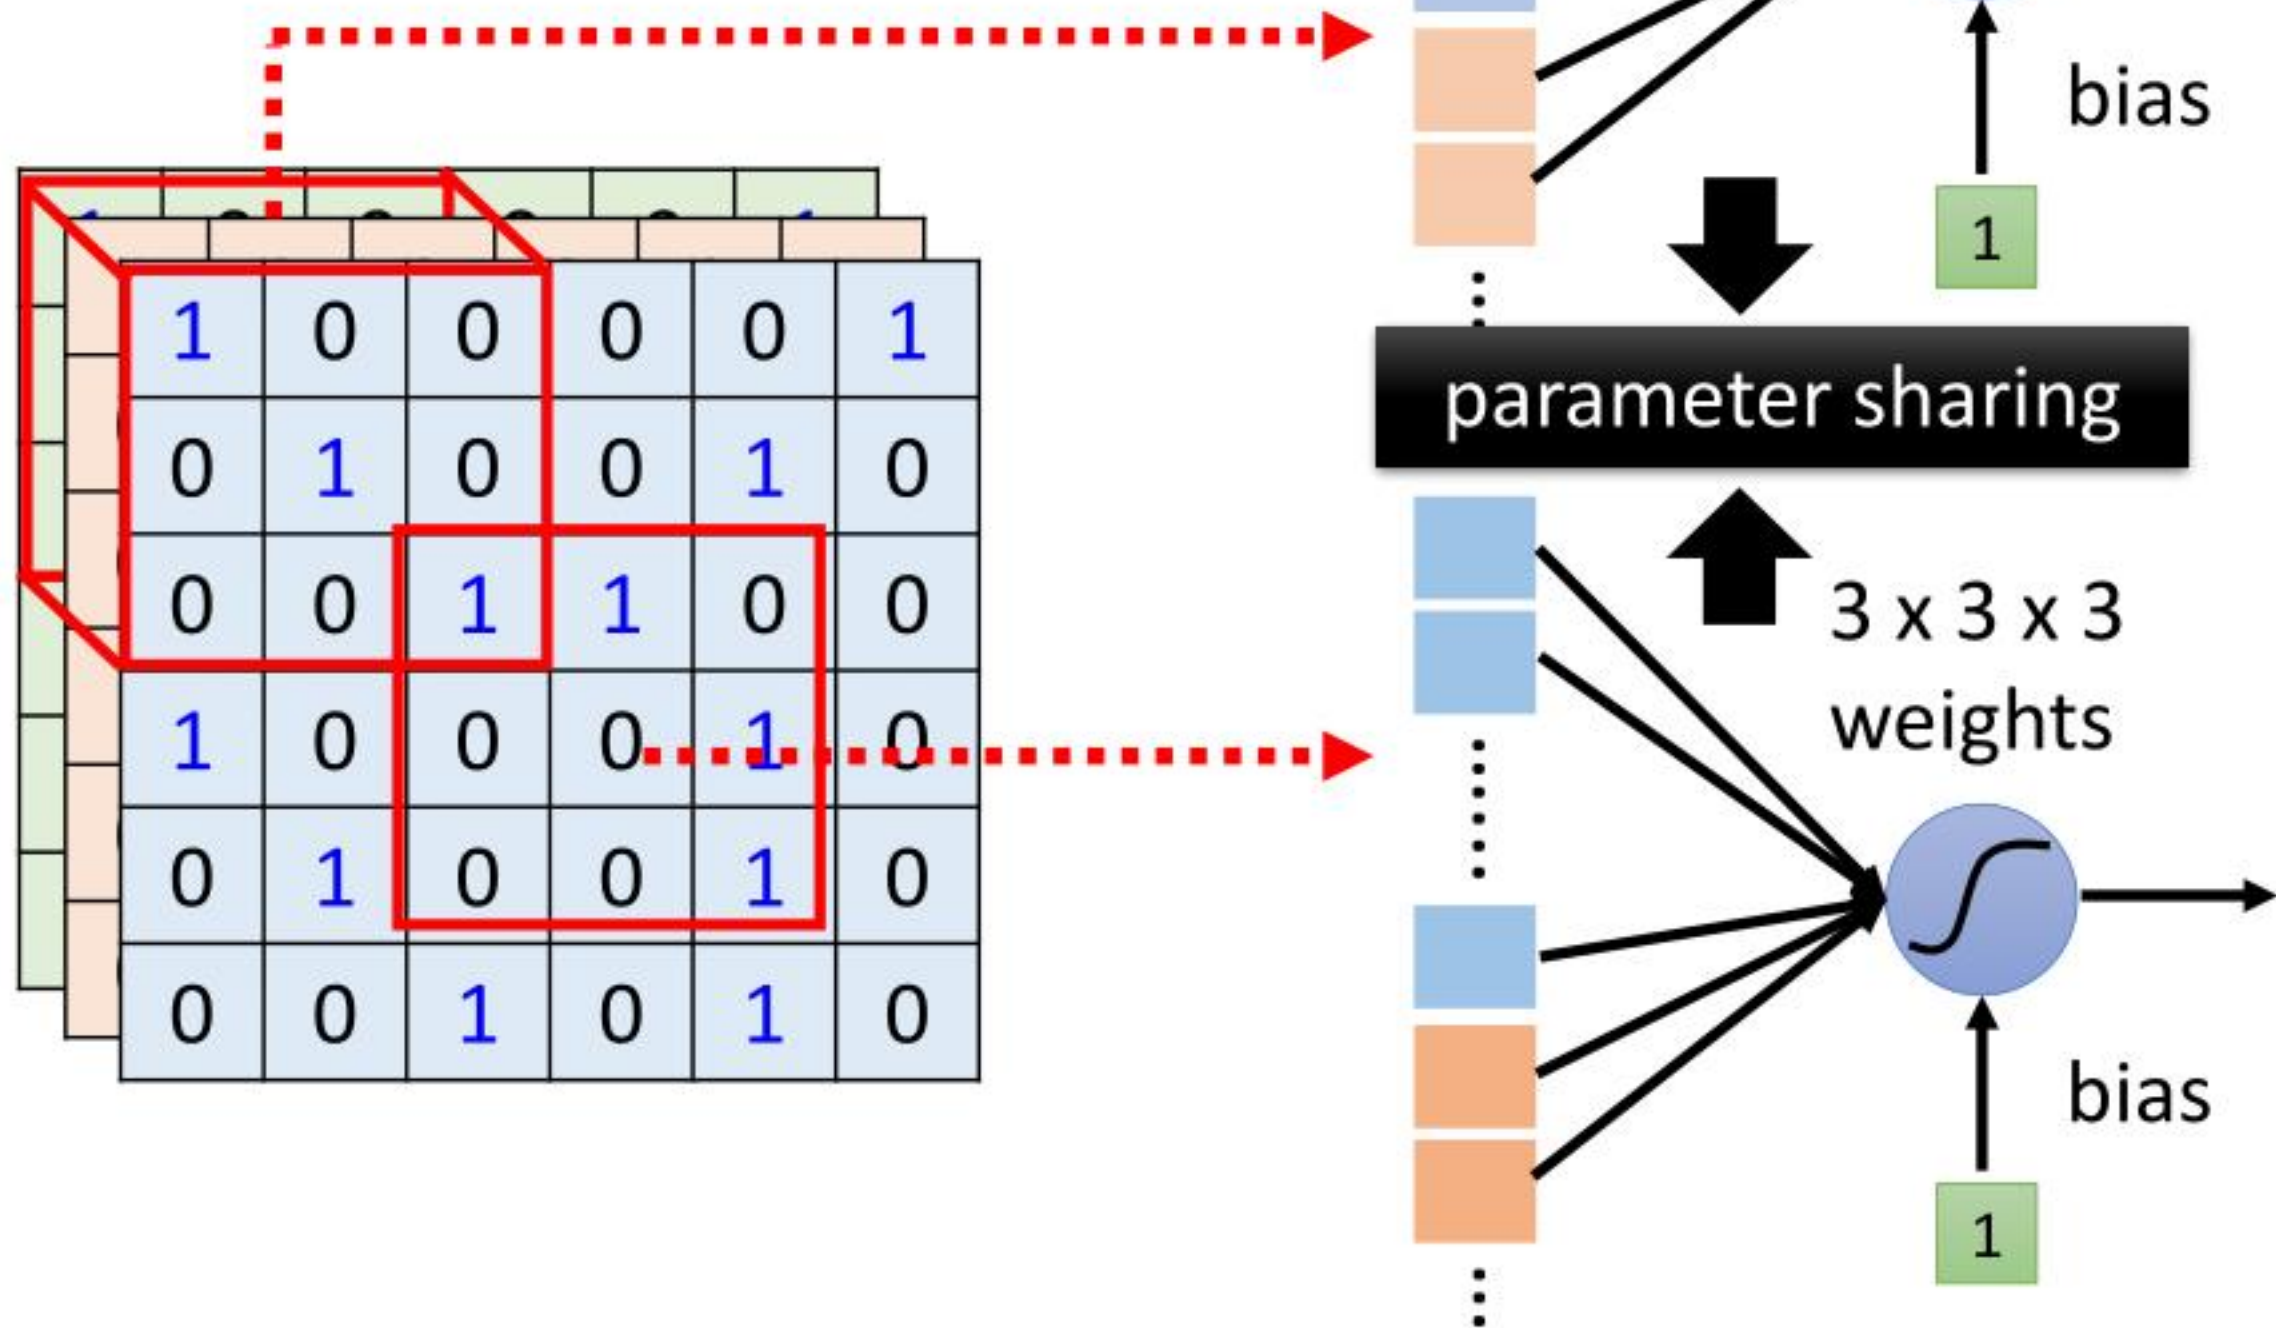


Observation 2


- The same patterns appear in different regions.




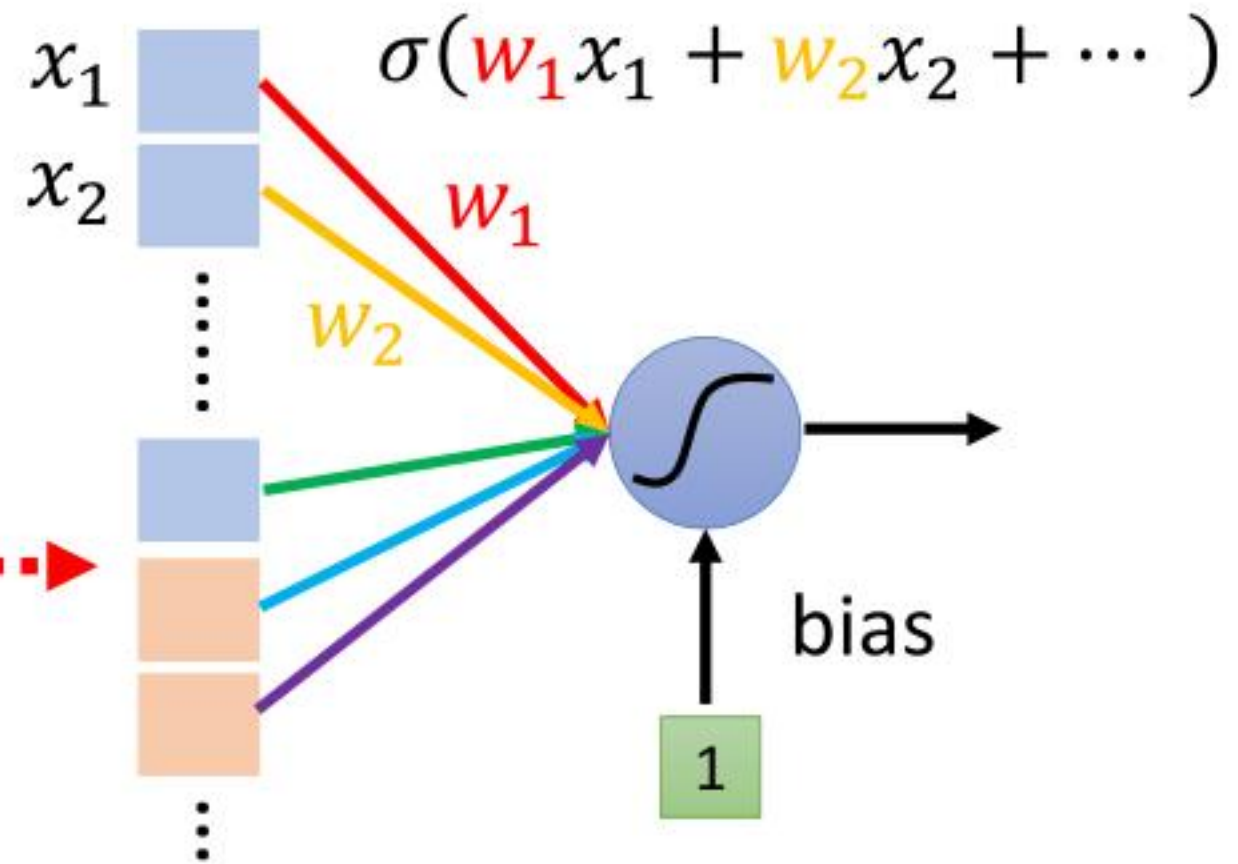
Simplification 2



Simplification 2




Two neurons with the same receptive field would not share parameters.



Simplification 2 – Typical Setting


Each receptive field has a set of neurons (e.g., 64 neurons).



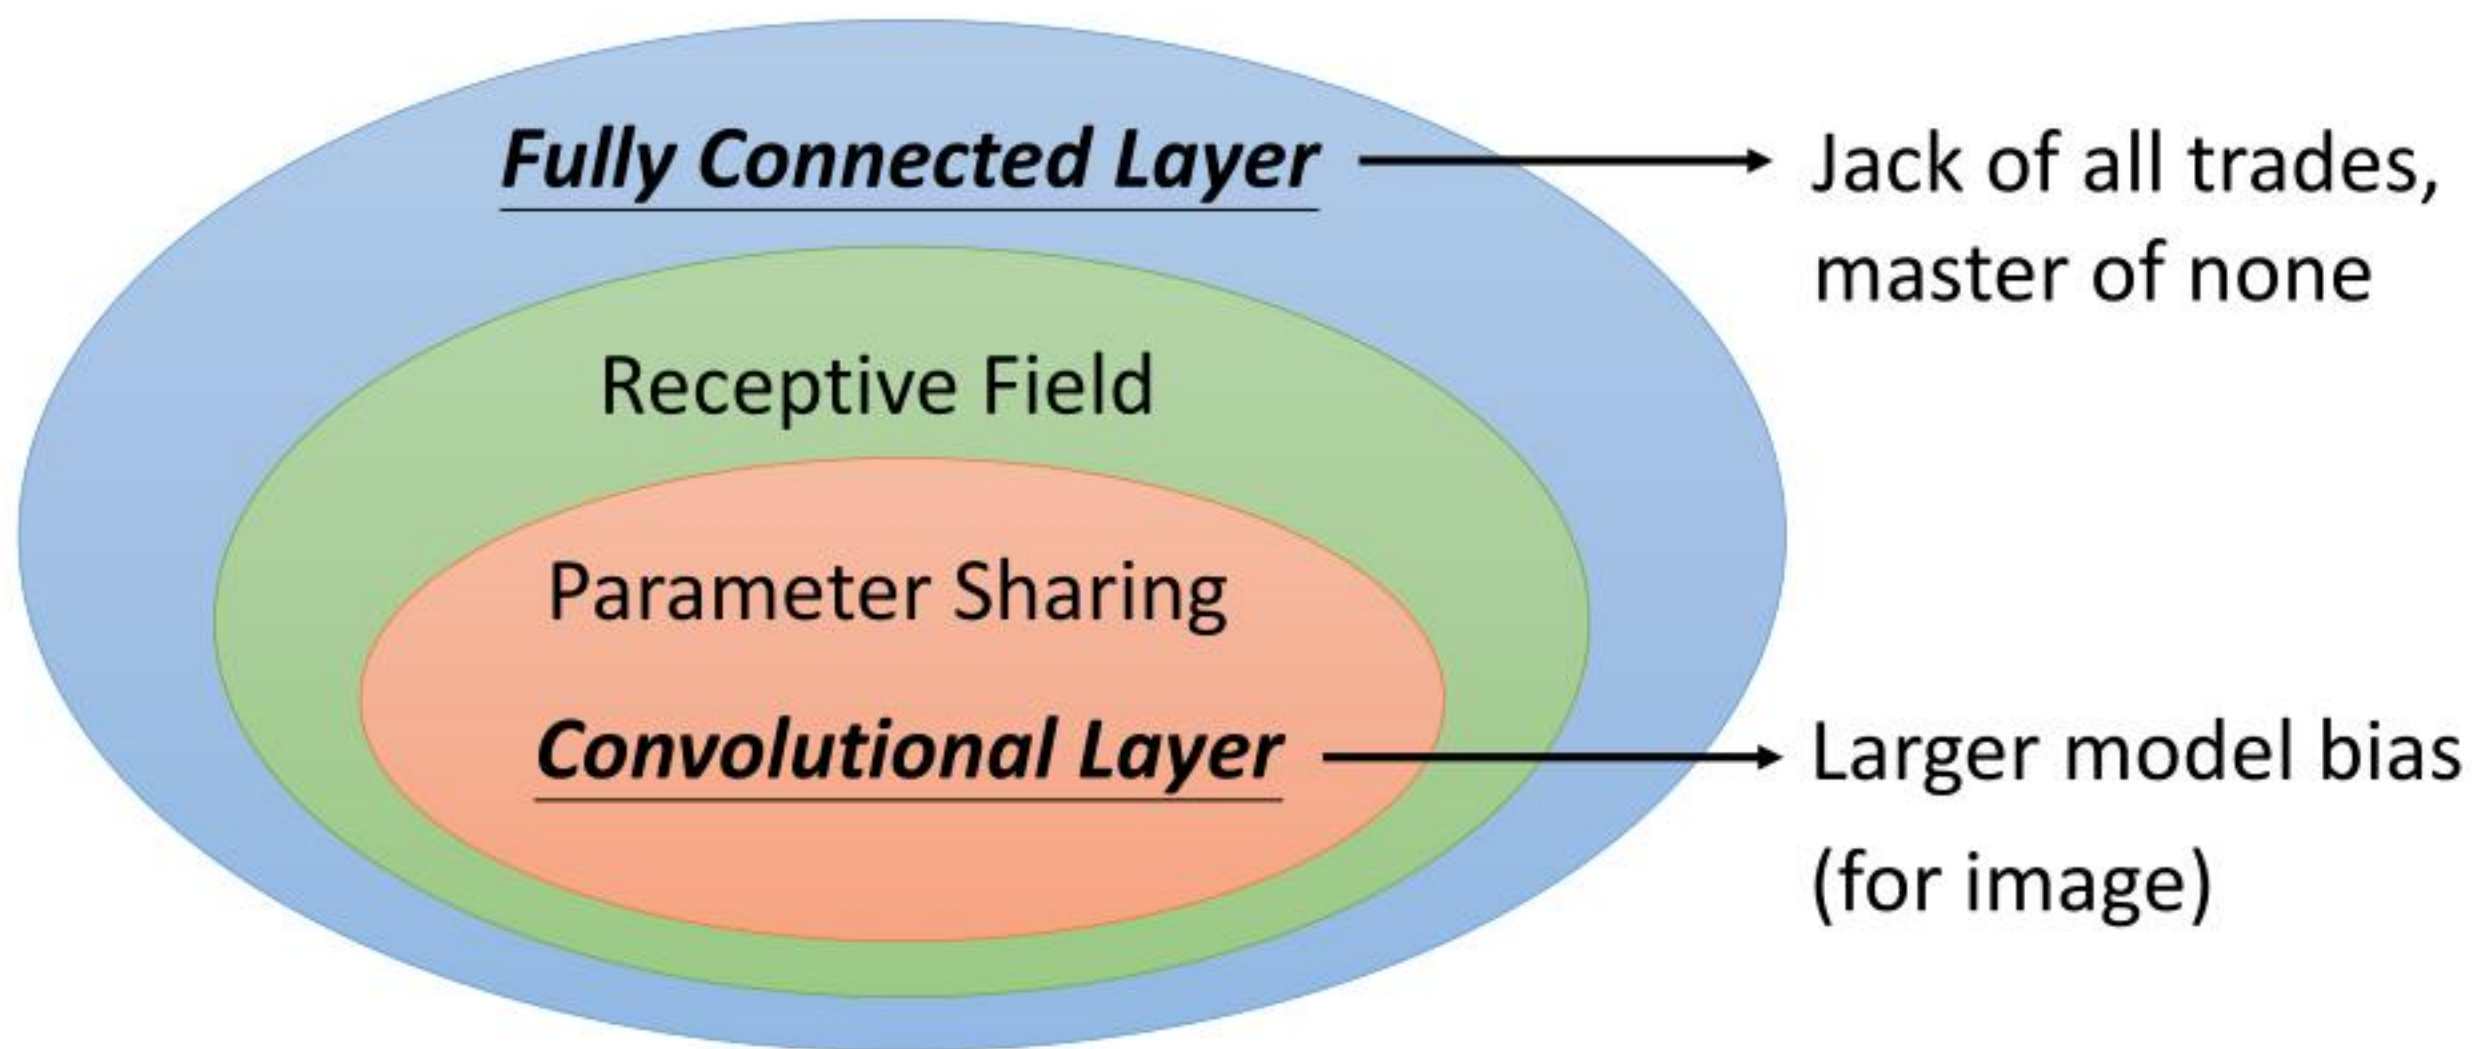
Simplification 2 – Typical Setting

Each receptive field has a set of neurons (e.g., 64 neurons).

Each receptive field has the neurons with the same set of parameters.



Benefit of Convolutional Layer



- Some patterns are much smaller than the whole image.
- The same patterns appear in different regions.

Another story based on *filter* 😊

Convolutional Layer




Convolution




⋮

channel = 3 (colorful)

channel = 1 (black and white)



Filter 1
3 x 3 x channel
tensor



Filter 2
3 x 3 x channel
tensor

⋮

Each filter detects a small pattern (3 x 3 x channel).

CNN – Convolution

Those are the network parameters to be learned.

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

6 x 6 image

1	-1	-1
-1	1	-1
-1	-1	1

Filter 1
Matrix

-1	1	-1
-1	1	-1
-1	1	-1

Filter 2
Matrix

⋮

Property 1

Each filter detects a small pattern (3 x 3).

CNN – Convolution


stride=1

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

6 x 6 image

1	-1	-1
-1	1	-1
-1	-1	1

Filter 1



CNN – Convolution


1	-1	-1
-1	1	-1
-1	-1	1

Filter 1

If stride=2

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

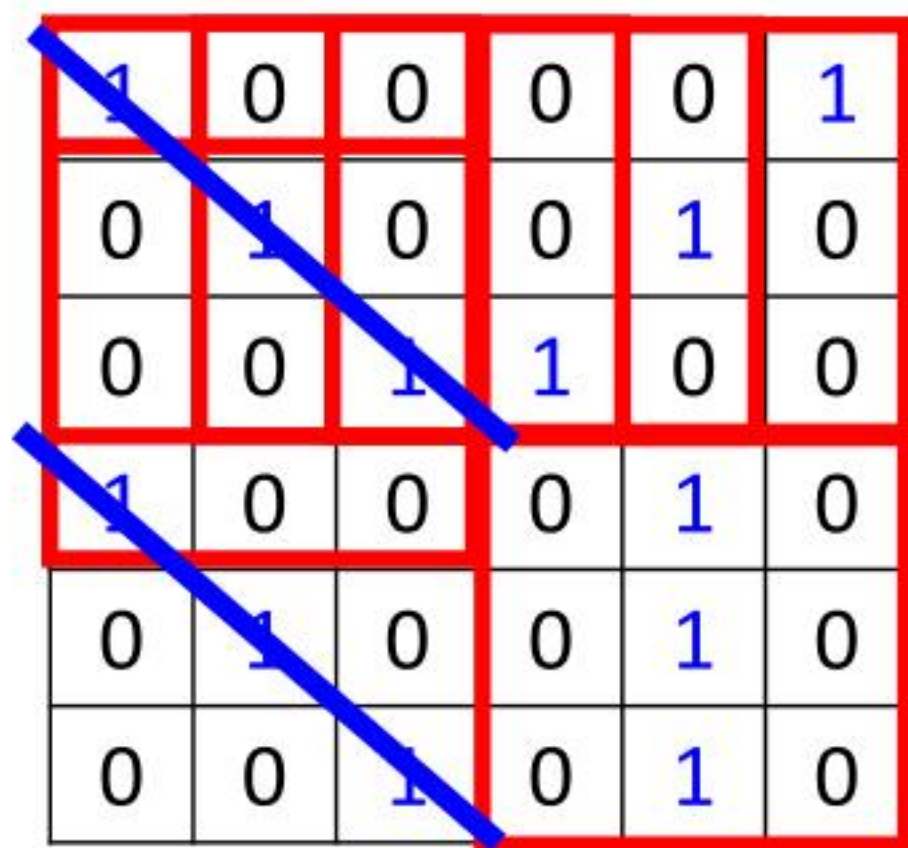
6 x 6 image



We set stride=1 below

CNN – Convolution

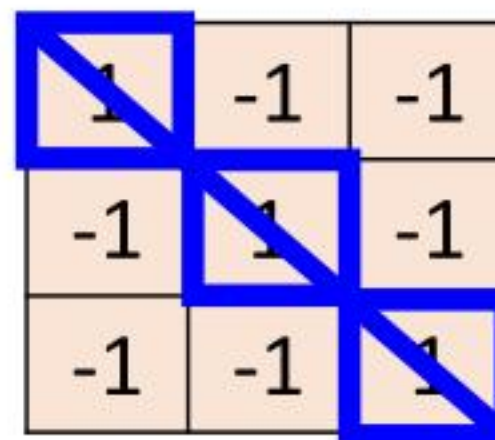
stride=1



A 6x6 grid of numbers representing an image. The grid is outlined with a red border. A 3x3 sub-region in the top-left corner is highlighted with a red border. A blue diagonal line runs from the top-left to the bottom-right of the entire grid.

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

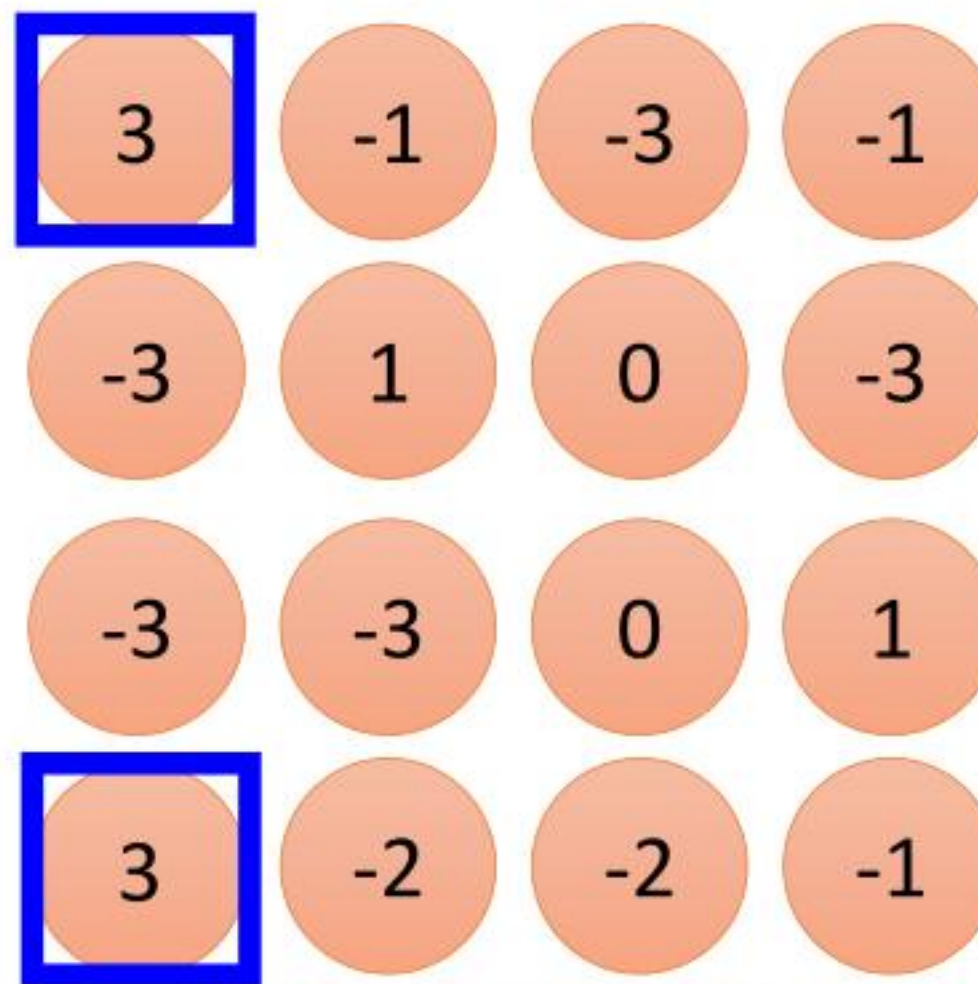
6 x 6 image



A 3x3 grid of numbers representing a filter. The grid is outlined with a blue border. A blue diagonal line runs from the top-left to the bottom-right of the grid.

1	-1	-1
-1	1	-1
-1	-1	1

Filter 1



A 4x4 grid of orange circles, each containing a number. The top-left circle (3) and the bottom-left circle (3) are highlighted with blue square borders.

3	-1	-3	-1
-3	1	0	-3
-3	-3	0	1
3	-2	-2	-1

Property 2

Convolutional Layer

-1	1	-1
-1	1	-1
-1	1	-1


Filter 2

stride=1

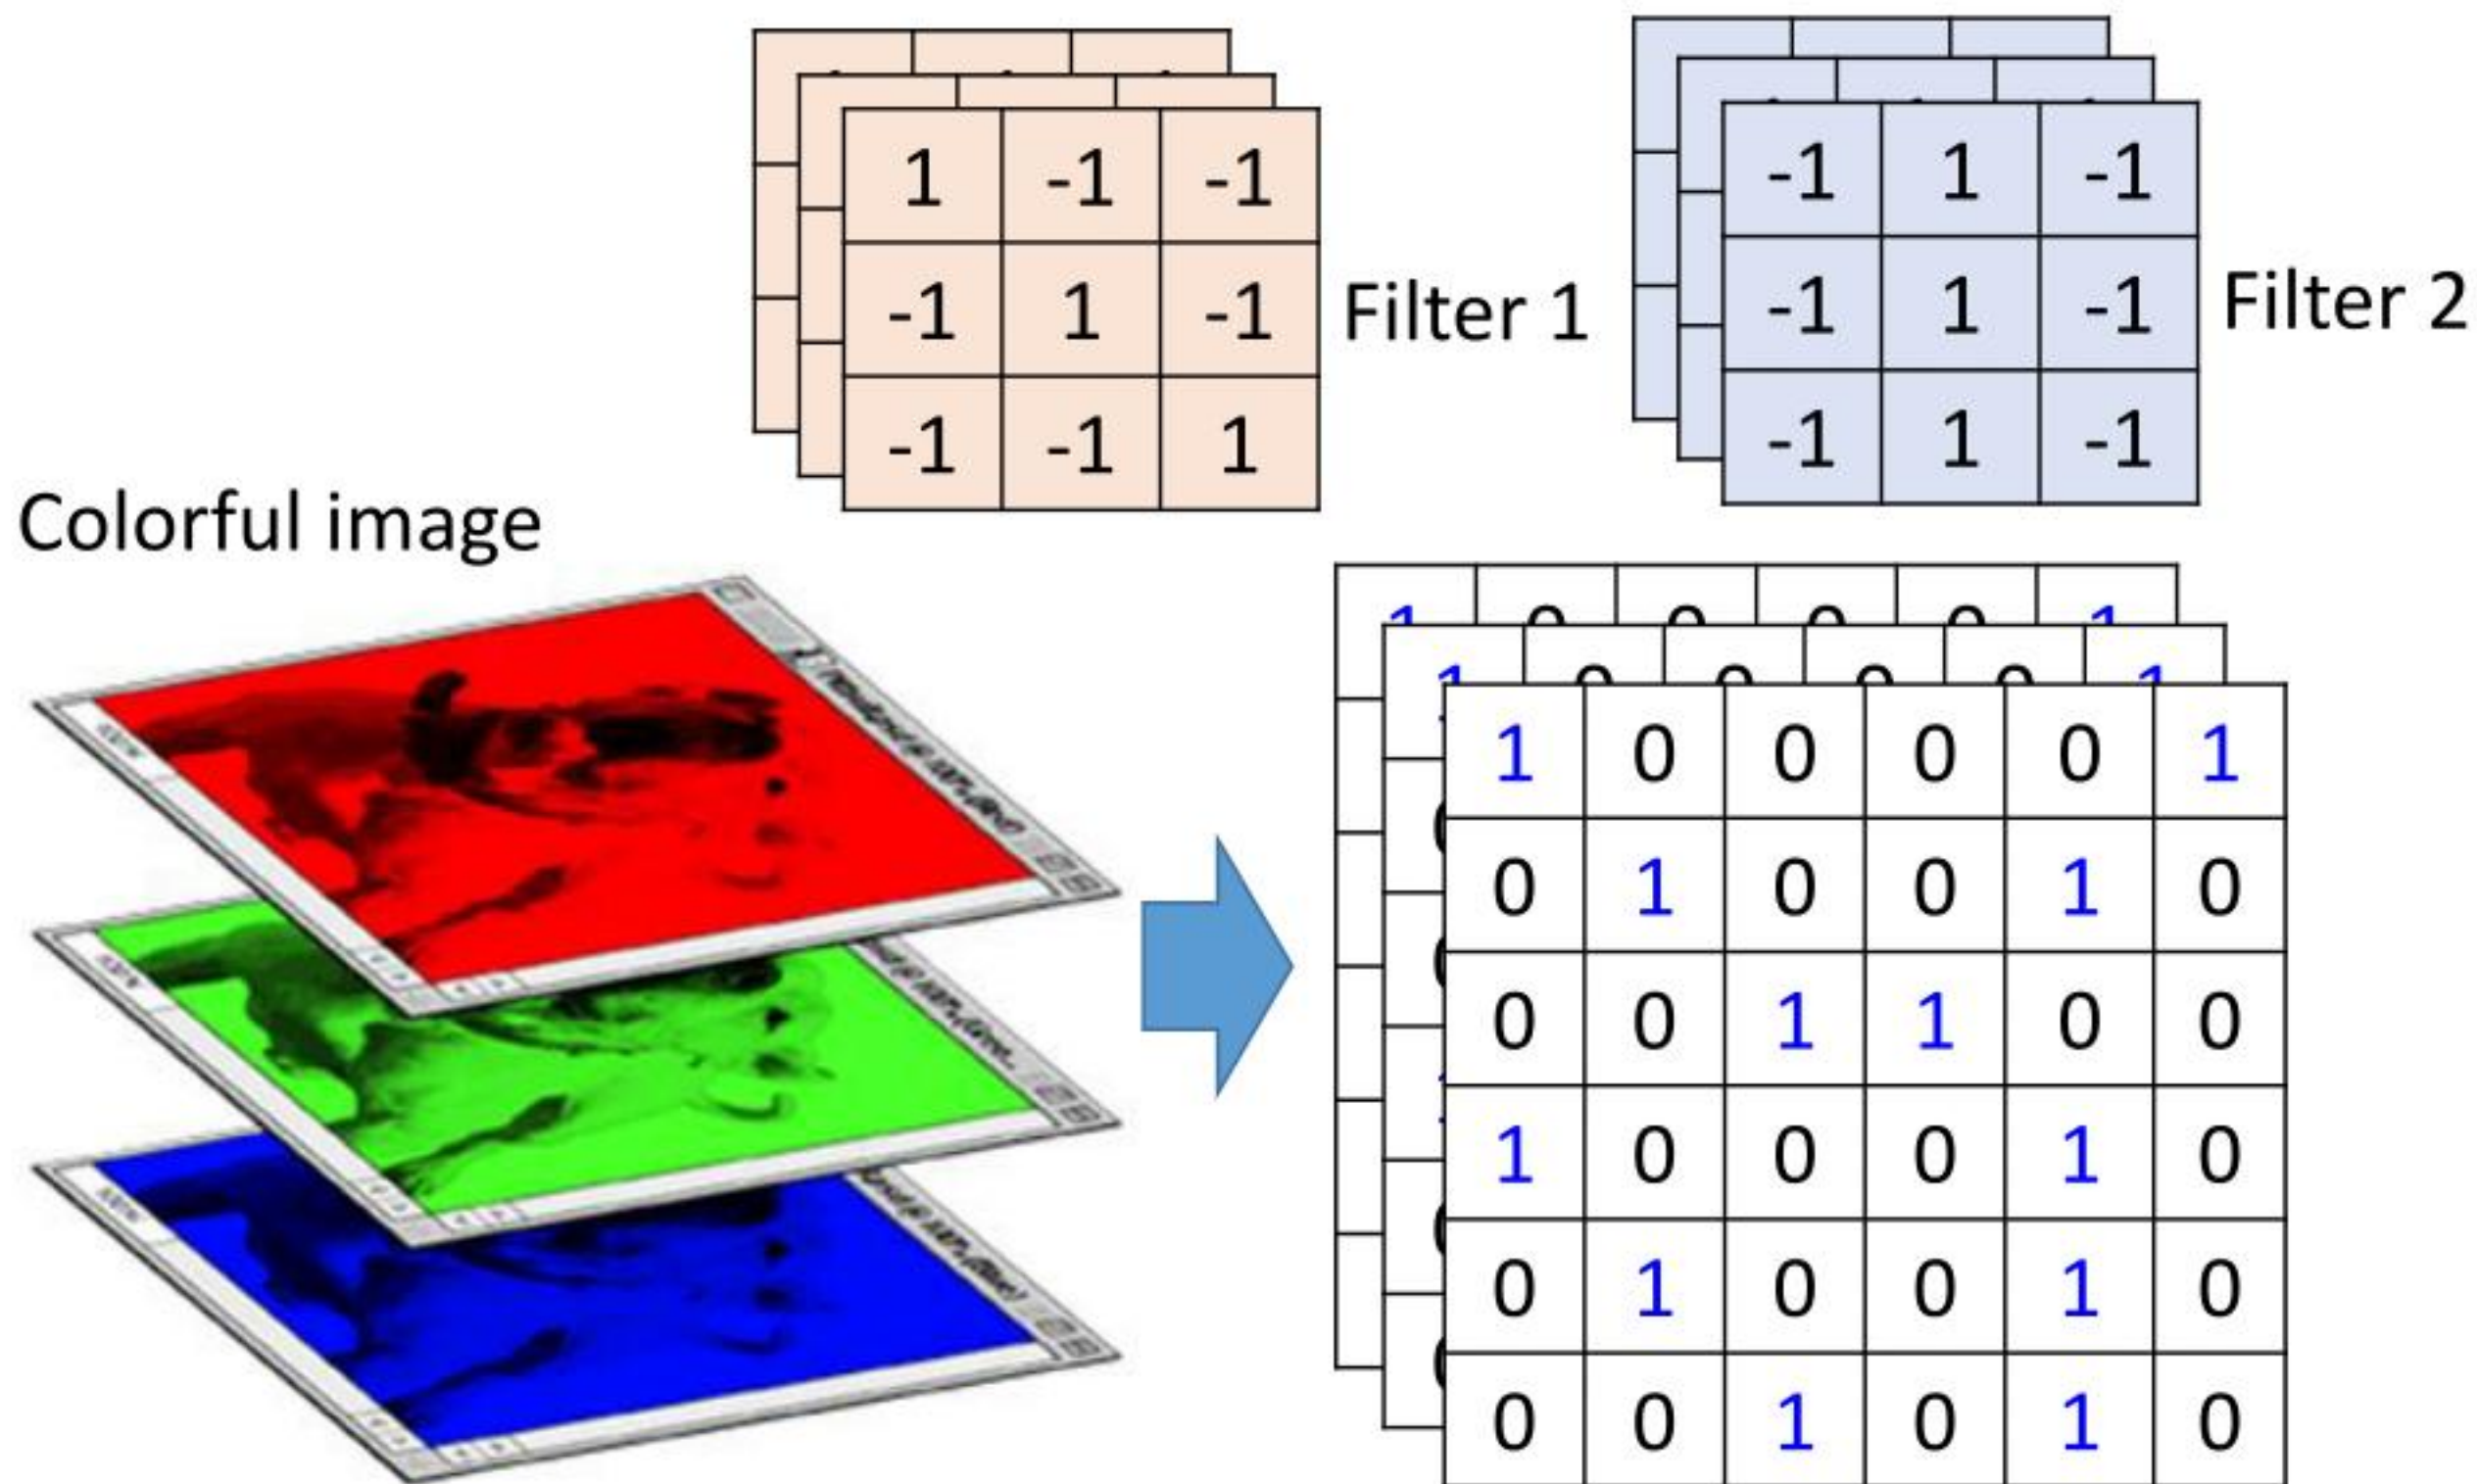
1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0

6 x 6 image

Do the same process for every filter



CNN – Colorful image



Convolutional Layer




64
filters

Convolution

Convolution

⋮



"Image" with 64 channels


Multiple Convolutional Layers




64
filters

Convolution

Convolution



"Image" with 64 channels



Multiple Convolutional Layers




64
filters

Convolution


Convolution

⋮

1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0






Convolution v.s. Fully Connected




Fully-
connected


1	0	0	0	0	1
0	1	0	0	1	0
0	0	1	1	0	0
1	0	0	0	1	0
0	1	0	0	1	0
0	0	1	0	1	0






Less parameters!






6 x 6 image




Less parameters!


Even less parameters!




Comparison of Two Stories



The neurons with different receptive fields **share the parameters**.



Each filter convolves over the input image.



Convolutional Layer

<u><i>Neuron Version Story</i></u>	<u><i>Filter Version Story</i></u>
Each neuron only considers a receptive field.	There are a set of filters detecting small patterns.
The neurons with different receptive fields share the parameters.	Each filter convolves over the input image.

They are the same story.

Observation 3

- Subsampling the pixels will not change the object

bird



subsampling

bird



We can subsample the pixels to make image smaller



Less parameters for the network to process the image

Pooling – Max Pooling

1	-1	-1
-1	1	-1
-1	-1	1

Filter 1

-1	1	-1
-1	1	-1
-1	1	-1

Filter 2

3	-1	-3	-1
-3	1	0	-3
-3	-3	0	1
3	-2	-2	-1

-1	-1	-1	-1
-1	-1	-2	1
-1	-1	-2	1
-1	0	-4	3


Convolutional Layers + Pooling




Convolution

Pooling

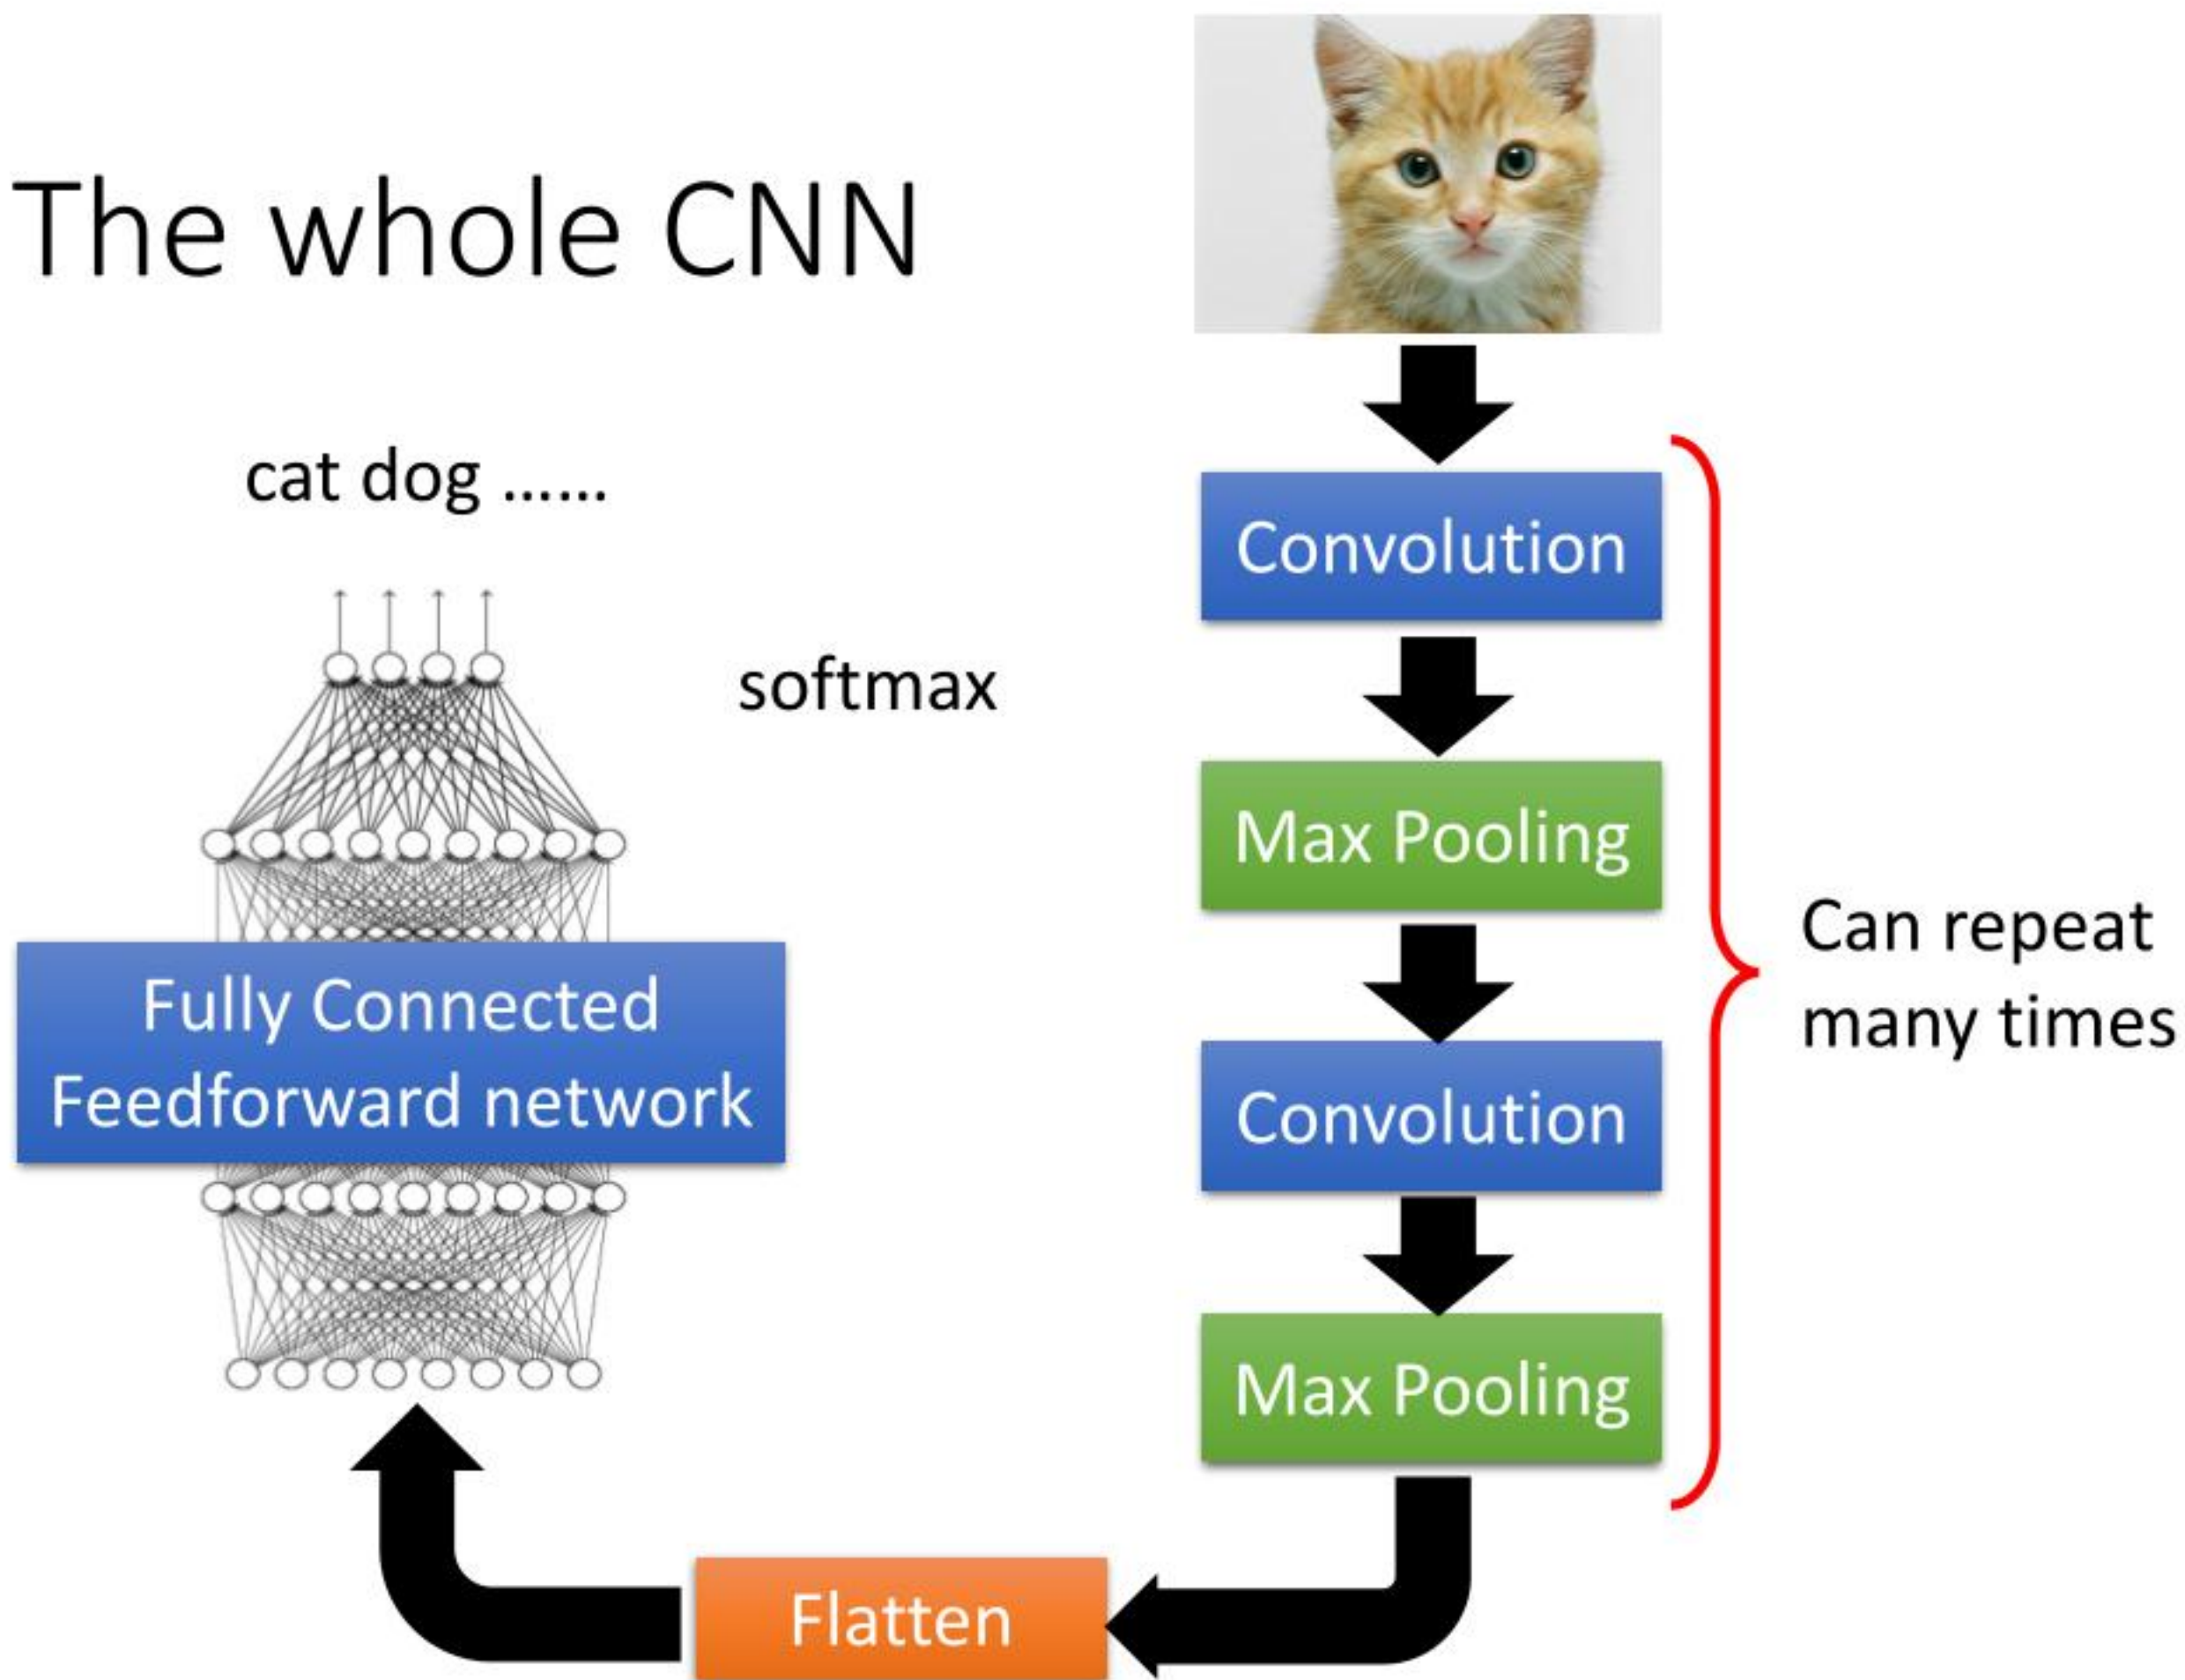
Repeat



"Image" with 64 channels



The whole CNN



The whole CNN

Property 1

- Some patterns are much smaller than the whole image

Property 2

- The same patterns appear in different regions.

Property 3

- Subsampling the pixels will not change the object



Convolution

Max Pooling

Convolution


Max Pooling

Flatten

Can repeat many times

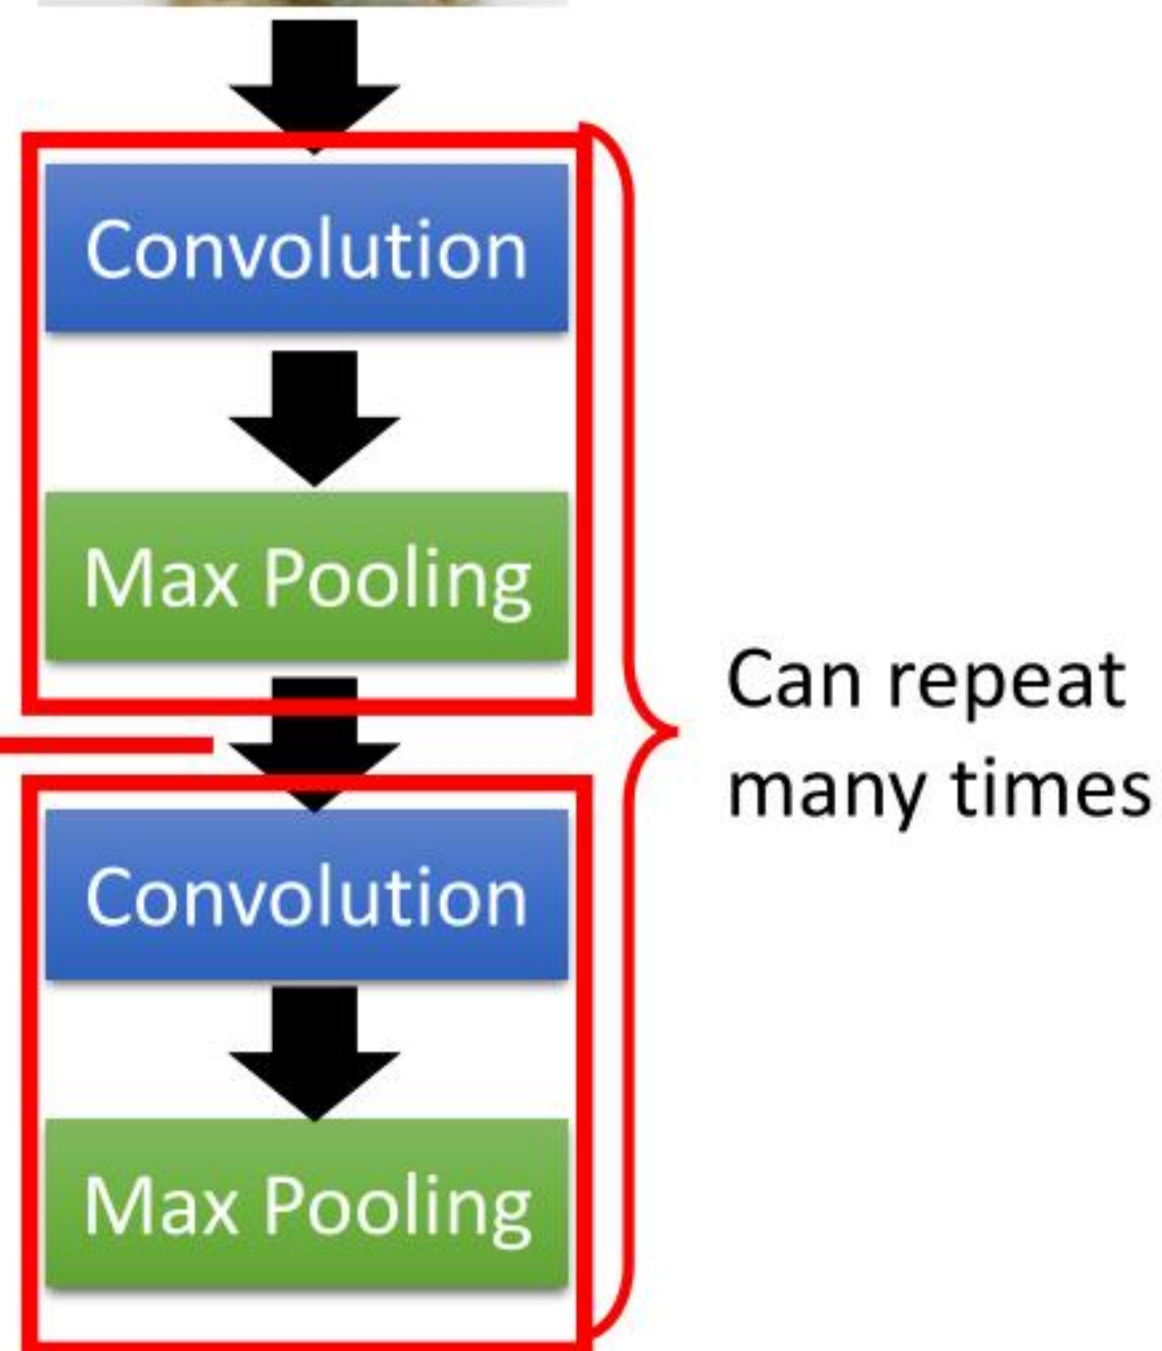


The whole CNN



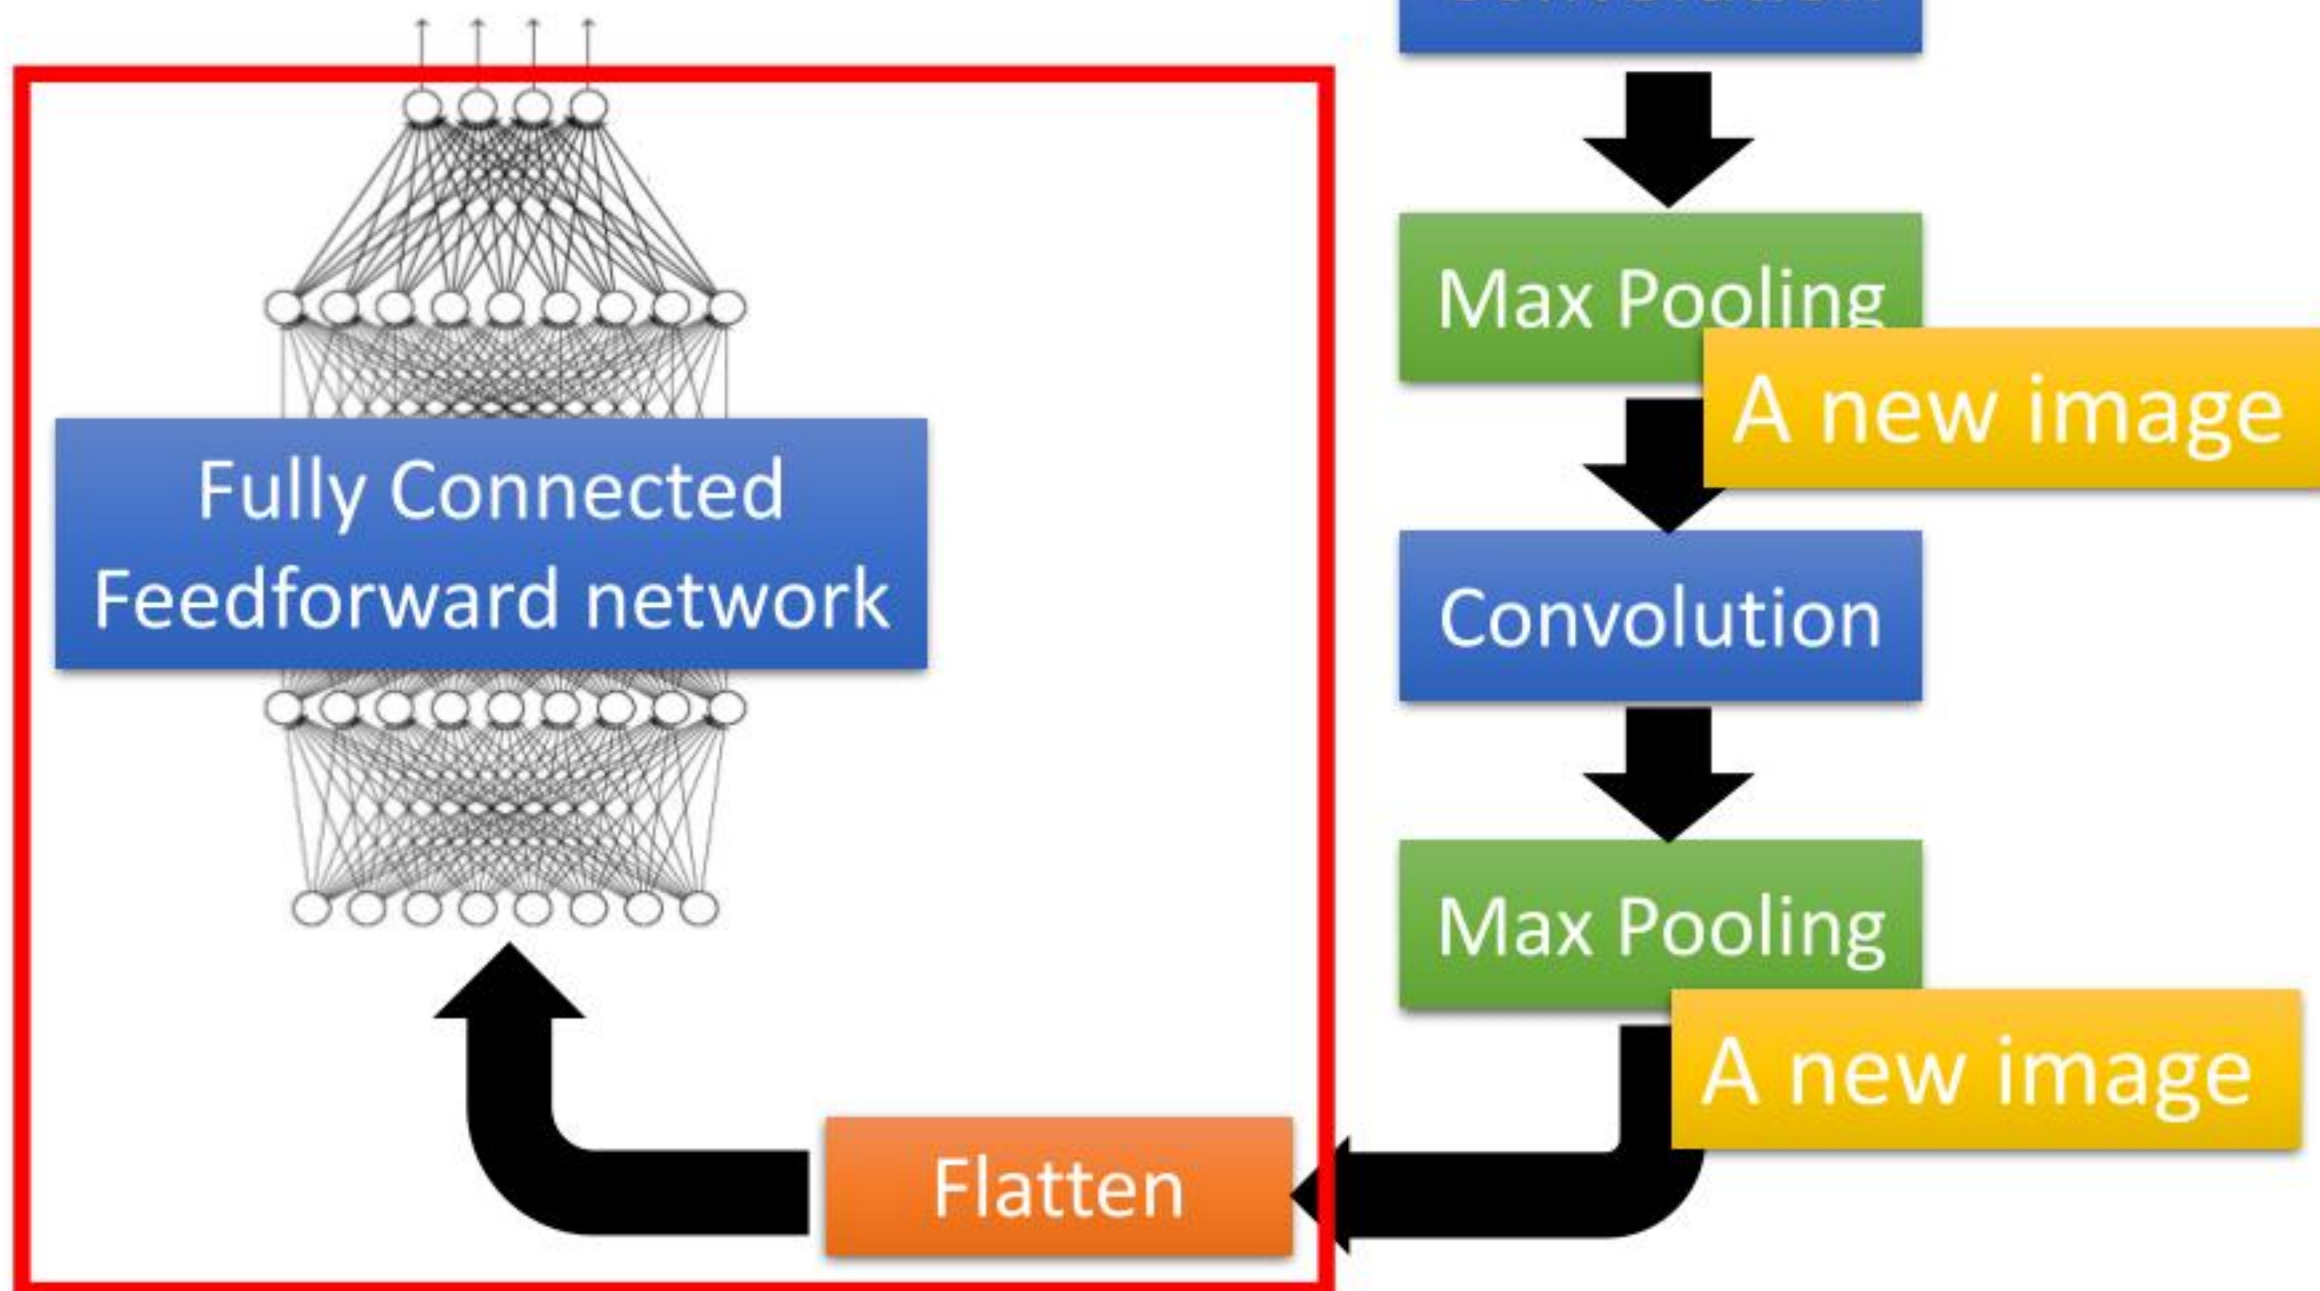
Smaller than the original image

The number of the channel is the number of filters

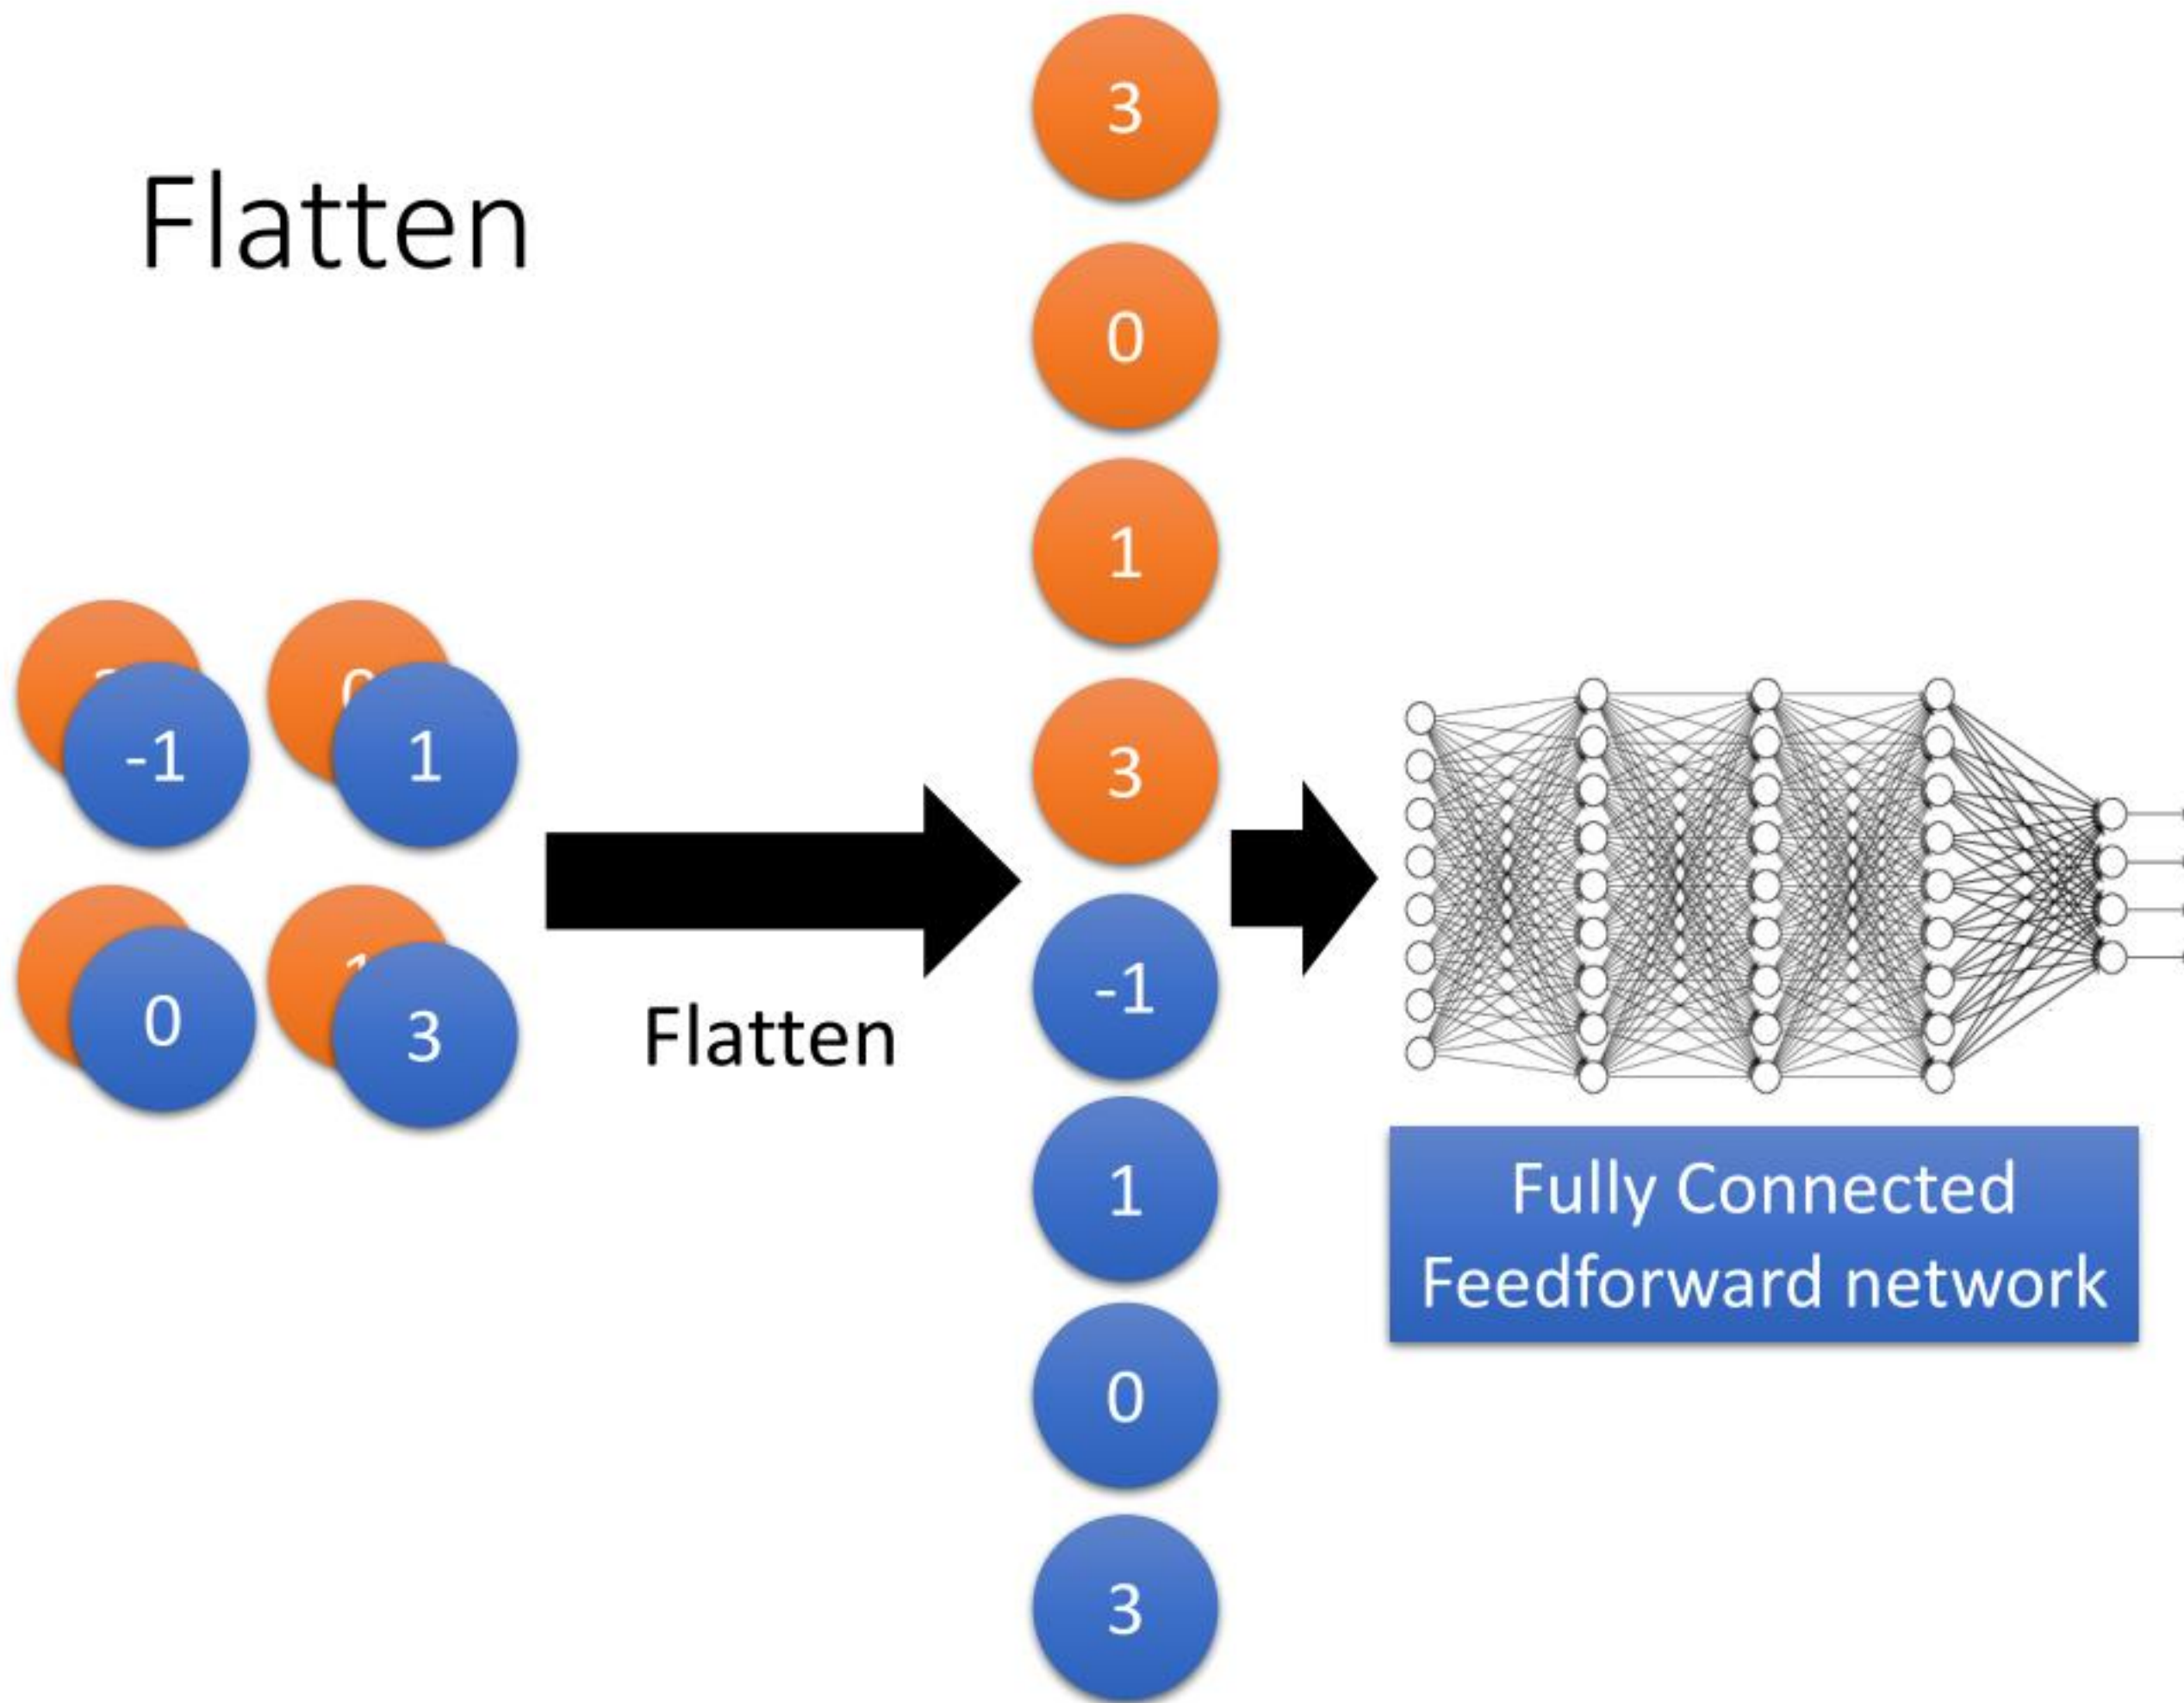


The whole CNN

cat dog



Flatten



Application: Playing Go



48 channels
in Alpha Go

Black: 1
white: -1
none: 0



Next move
(19 x 19
positions)
19 x 19 classes

Fully-connected
network can be used

But CNN performs much better.



Why CNN for Go playing?

- Some patterns are much smaller than the whole image

Alpha Go uses 5 x 5 for first layer



- The same patterns appear in different regions.



Why CNN for Go playing?

- Subsampling the pixels will not change the object



Pooling

How to explain this???

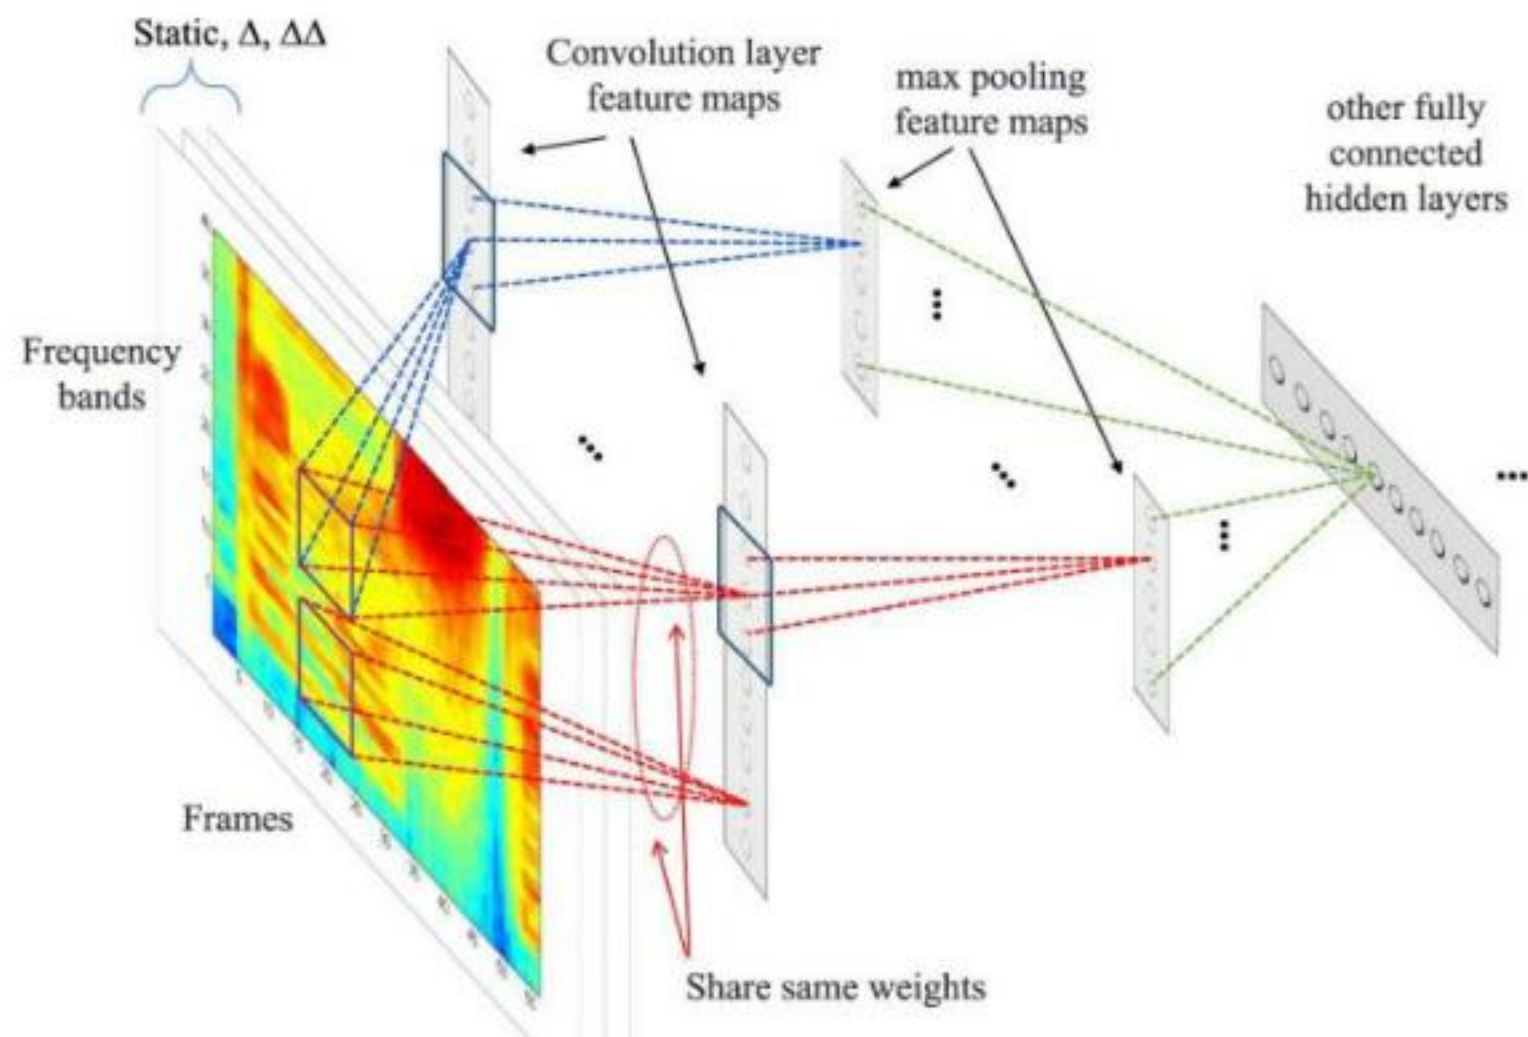
Neural network architecture. The input to the policy network is a $19 \times 19 \times 48$ image stack consisting of 48 feature planes. The first hidden layer zero pads the input into a 23×23 image, then convolves k filters of kernel size 5×5 with stride 1 with the input image and applies a rectifier nonlinearity. Each of the subsequent hidden layers 2 to 12 zero pads the respective previous hidden layer into a 21×21 image, then convolves k filters of kernel size 3×3 with stride 1, again followed by a rectifier nonlinearity. The final layer convolves 1 filter of kernel size 1×1 with stride 1, with a different bias for each position, and applies a softmax function. The match version of AlphaGo used $k = 192$ filters; Fig. 2b and Extended Data Table 1. The match version of AlphaGo used 256 and 384 filters.

Alpha Go does not use Pooling

More Applications

Speech

<https://dl.acm.org/doi/10.1109/TASLP.2014.2339736>



Natural Language Processing

<https://www.aclweb.org/anthology/S15-2079/>

